

Abstract

Null hypothesis significance testing (NHST) is cited as a threat to validity and reproducibility. While many individuals suggest we focus on altering the p -value at which we deem an effect significant, we believe this suggestion is short-sighted. Alternative procedures (i.e., Bayesian analyses and Observation Oriented Modeling: OOM) can be more powerful and meaningful to our discipline. However, these methodologies are less frequently utilized and are rarely discussed in combination with NHST. Herein, we discuss three methodologies (NHST, Bayesian Model comparison, and OOM), then compare the possible interpretations of three analyses (ANOVA, Bayes Factor, and an Ordinal Pattern Analysis) in various data environments using a frequentist simulation study. We found that changing significance thresholds had little effect on conclusions. Further, we suggest that evaluating multiple estimates as evidence of an effect allows for more robust and nuanced interpretations of results and implies the need to redefine evidentiary value and reporting practices.

Keywords: null hypothesis testing, p -values, Bayes Factors, Observation Oriented Modeling, evidence

Beyond p -values: Utilizing Multiple Methods to Evaluate Evidence

Recent events in psychological science have prompted concerns within the discipline regarding research practices and ultimately, the validity and reproducibility of published reports (Etz & Vandekerckhove, 2016; Lindsay, 2015; Open Science Collaboration, 2015; van Elk et al., 2015). One often discussed matter is over-reliance, abuse, and potential hacking of p -values produced by frequentist null hypothesis significance testing (NHST), as well as misinterpretations of NHST results (Gigerenzer, 2004; Ioannidis, 2005; Simmons, Nelson, & Simonsohn, 2011). We agree with these concerns and believe that many before us have voiced sound, generally accepted opinions on potential remedies, such as an increased focus on effect sizes (Cumming, 2008; Lakens, 2013; Maxwell, Lau, & Howard, 2015; Nosek, Spies, & Motyl, 2012).

However, other suggestions have been met with less enthusiasm, including an article by Benjamin et al. (2018) advocating that researchers should begin thinking only of p -values less than .005 as “statistically significant”, thus changing α levels to control Type I error rates. Alternatively, Pericchi and Pereira (2016) promote the use of fluctuating α levels as a function of sample size to assist with these errors. Trafimow et al. (2018) critiques this suggestion to broadly lower the α level to .005 and suggested that findings should be weighted on the basis of evidence accumulation from multiple studies. We argue that α should not be the sole focus of our attention, but rather if a p -value should be utilized at all, and, if so, what that p -value can tell you in relation to other indicators. While NHST and p -values may have merit, researchers have a wealth of other statistical tools available to them. We believe that improvements may be made to the sciences as a whole when individuals become aware of the these tools and how these methods may be used, either alone or in combination, to strengthen understanding of data and conclusions. These sentiments have been shared by the American Statistical Association who recently held a conference focusing on going beyond NHST, expanding their previous stance p -values

43 (Wasserstein & Lazar, 2016).

44 Therefore, the main goal of this project was to show researchers how two alternative
45 paradigms compare to NHST in terms of their methodological design, statistical
46 interpretations, and comparative robustness. Herein, we will discuss the following
47 methodologies: NHST, Bayes Factor comparisons, and Observation Oriented Modeling. In
48 order to compare their methodological designs, we first provide historical backgrounds,
49 procedural steps, and limitations for each paradigm. We then simulated data using a three
50 time point repeated measures design with a Likert-type scale as the outcome variable to be
51 able to compare the statistical interpretations and comparative robustness. By simulating
52 possible datasets and analyzing them with each of the three paradigms, we will be able to
53 discuss the conclusions these three methods reach given the same data and to compare how
54 often these methodologies agree within different data environments (i.e., given varying
55 sample sizes and effect sizes). Beyond simply comparing methodologies, we also sought to
56 identify how changing the α criteria within the NHST framework may alter conclusions.
57 Although previous work has already compared Frequentist NHST to Bayesian approaches
58 (Goodman, 1999; Rouder, Morey, Speckman, & Province, 2012; Wetzels et al., 2011), this
59 manuscript adds a novel contribution: Observation Oriented Modeling. By introducing social
60 scientists to Observation Oriented Modeling (OOM), a relatively new paradigm that is
61 readily interpretable, we will show both how useful this paradigm can be in these contexts,
62 and how it compares to two well-known methods. We hope that by discussing these
63 methodologies in terms of a simple statistical analysis researchers will be able to easily
64 compare and contrast methodologies.

Null Hypothesis Significance Testing

History

Many attribute the frequentist NHST procedure to Ronald A. Fisher (Fisher, 1932). However, Fisher's ideas are a far cry from the NHST procedure implemented today. Fisher believed in creating one "null" hypothesis, which he described as a hypothesis to be "nullified", or shown incorrect, not as a zero-difference hypothesis (Lehmann, 2011). He also believed that the use of any omnibus level of significance showed a "lack of statistical thinking" (Gigerenzer et al., 2004). He instead believed we should report the exact significance value of a test and let others make their own decision about the claims, which is more in line with the typical reporting recommendations provided by the American Psychological Association (American Psychological Association, 2010). Fisher spoke of this work to William Gosset, the man who created the Student's t -test and contributed work on the correlation coefficient (Lehmann, 2011). Gosset in turn discussed the idea of an alternative hypothesis, a piece not included in Fisher's procedure, with decision theorist Egon Pearson.

From this discussion, Egon Pearson and Jerzy Neyman created Neyman-Pearson decision theory. This theory consists of two hypotheses (i.e., null and alternative) and a binary decision criteria (i.e., significant or not, Lehmann, 1993). However, this combination created the possibility of researcher decision errors (Dienes, 2008). A researcher may falsely reject the null hypothesis (Probability of Type I error, α) or falsely fail to reject the null (Probability of Type II error, β). α levels set the binary decision criteria, which are used as the critical p -value for hypothesis testing (i.e., $p < .05$), and are thus seen as evidence to reject the null hypothesis. β and power are inherently linked (Power = $1 - \beta$), so as the likelihood of finding a true effect increases beta decreases (Maxwell & Delaney, 2004). Although α values can be chosen to be quite small, and methods (such as decreasing error

variance or using a one-tailed test in contrast to a two-tailed test) can decrease β values as well, a researcher can never know if they have made the correct decision, or a decision error. Thus, Neyman and Pearson clearly state that a hypothesis should not be blindly supported based solely on the estimates of one statistical test, and that replication and reproduction of results are imperative. The recent work of the Open Science Collaboration (2015) has also highlighted the need for replication studies and interpretation of results in an appropriate context. Additionally, Neyman and Pearson emphasized that use of set α s and β s is illogical and sought instead for researchers to adjust their analysis to the needs of the particular task at hand (Gigerenzer, 2004).

Typical Procedure

Neither Fisher’s hypothesis testing, nor Neyman-Pearson decision theory quite match the NSHT procedure as it is taught and applied today. Psychologists have largely adopted an amalgamation of the two approaches. Here, we attempt to outline what we believe is the most appropriate way to carry out the traditional NHST procedure in the context of a repeated measures ANOVA with three levels, although we note that this set of steps is not necessarily how researchers carry out the procedure in practice:

- 1) Create two hypotheses, one to be “nullified” and one “alternative” hypothesis. Within this repeated measures framework, most researchers would define a null hypothesis (H_0) that indicates of all three time point population means (μ) are equal (i.e., all of our observed values X_i , regardless of which time point they were assessed at X_{ij} , arise from a normal distribution N with some mean μ and variance σ^2). The alternative hypothesis (H_A) would then be that each mean (μ) is allowed to be different from the grand mean by some amount (δ ; as we now have observations being drawn from three potential normal distributions, all of which may have a different mean value, but the same variance). Within this frequentist framework, the observed data X_i are

considered the expression of random variables, and the parameter μ is considered to be fixed but unknown. These can be operationalized as follows:

$$H_0 : X_{ij} \sim N(\mu, \sigma^2)$$

$$H_A : X_{ij} \sim N(\mu + \delta_i, \sigma^2)$$

- 2) Select an α level that is appropriate given the context of your research, your analysis plan, and your research question, and do not blindly adopt an omnibus critical p -value (Lakens et al., 2018; Lehmann, 2011). Again, we reiterate that such α justification and selection is not necessarily how all researchers approach these tests.
- 3) Compute your given analysis and identify the corresponding p -value. If your p -value is less than the chosen α , reject the null hypothesis and state that there appear to be differences between some of your population means; however, if your p -value is greater than or equal to the value selected, do not reject the null hypothesis, and state that a difference between the population means could not be supported.

The specific analysis used to test these models here, the repeated measures ANOVA with three levels, requires some additional assumptions that must be met before an analysis is begun (Tabachnick & Fidell, 2012). Data need to have no outlying or influential observations. Data must have a normal sampling distribution, be linearly related, and have independent errors. Depending on the statistical test, data must also be checked for equal variances, sphericity, and additivity. These assumptions can be checked and, if necessary, corrected for; however, violations of these assumptions can lead to inaccurate decisions and attenuated power. Further, with many analysis programs, data are required to have no missing values.

While this approach is widely used, there are many limitations associated with it. First, this method can be sensitive to violations of the stated assumptions, and especially, if

the sample size is not large enough for the sampling distribution to approximate a normal distribution (Tabachnick & Fidell, 2012). Even if assumptions are met, or nonparametric tests are implemented (e.g., for situations where a normal distribution assumption cannot be met), this methodology does not allow a researcher to state anything about the absence of an effect (i.e., no true differences). The null hypothesis is assumed to be true and therefore, calculations are made to determine how likely it is that the data collected arose from that null distribution. Therefore, one can state that it is unlikely that sample(s) came from the null distribution (e.g., reject the null) but not the likelihood of the null hypothesis. Given the recent findings regarding reproducibility, showing support for the absence of an effect (e.g., supporting the null hypothesis) can be even more crucial than showing support for the presence of an effect (e.g., rejecting the null and supporting the alternative hypothesis; Bakker, van Dijk, & Wicherts, 2012; Lakens, 2017).

Bayes Factors

History

Thomas Bayes was a statistician and Presbyterian minister whose works are still influential today (Bellhouse, 2004). Bayes' theorem solved the inverse probability problem, namely that through the frequentist approach, one can only know the probability of data existing given a hypothesis being true, never the probability of a hypothesis being true given that the data exist (Dienes, 2008). Bayes' theorem allows one to calculate the probability of a hypothesis given some data (posterior belief) by using how probable one believes the hypothesis to be before data was collected (prior belief) and how probable one believes the data to be given one's hypothesis (likelihood). Thus, with his theorem, researchers are able to update (through the use of the likelihood) our initial beliefs (our prior) given some data (Gelman, Carlin, Stern, & Rubin, 2013). Pierre-Simon Laplace pioneered Bayesianism and advocated for a broader interpretation of this theorem (De Laplace, 1774). The use of

Bayesian statistics has been suggested as an NHST alternative (Dienes, 2014; Wagenmakers, 2007), but this approach has largely been undervalued in favor of frequentist methods as, until recently, Bayesian analysis required considerable computational effort. However, today we possess the technology necessary to efficiently conduct Bayesian analyses. While open source software, such as *R* and JASP, require minimal learning to be able to effectively operate (Morey & Rouder, 2015), researchers will need to invest more effort to understand the focus and interpretation of Bayes Factor (BF) comparisons as they differ from traditional NHST.

The Bayesian framework can be viewed as a continuum, with objective Bayesian analyses on one end, and subjective Bayesian analyses on the other (Press, 2002). While this topic could lend itself to its own manuscript, here we will simply summarize the two endpoints, and discuss where our analysis may be perceived to fall on the line. Objective Bayesian analysis is closest to frequentist theory, as the aim is to minimize the influence of priors through the use of non-informative priors (such as Jefferys priors that are designed to be invariant under reparameterization; Datta & Ghosh, 1996); thus, the data are allowed to maximally effect the posterior distribution. Further, objective Bayesian methods are influenced by the same quality criteria that frequentist methods used, including Type I error rate and power (Sellke, Bayarri, & Berger, 2001). On the other end, subjective Bayes analyses include rigorously informed priors so that current knowledge can play a large role in the posterior. Our current analysis splits these two; we do not utilize completely uniformed (objective) priors, as we can adjust for basic knowledge of the constraints of our data type.

It is worthwhile to note that we have discussed objective and subjective views of Bayesianism through the traditional lens (e.g., how uninformed the prior is). However, Berger (2006) has instead suggested that objective Bayesian approaches allow the priors to be driven by the data, while subjective Bayesian approaches allow the priors to be driven by the researcher. Given the usual lack of information about underlying distributions, a wider

band of inclusion was used for prior information for this study. The *BayesFactor* package (Morey & Rouder, 2015) assists greatly in the choice of prior and is especially user-friendly for applied researchers, as it makes use of recommended default priors that have been chosen to be safe assumptions under a broad range of data and topics (Rouder et al., 2012; Rouder, Speckman, Sun, Morey, & Iverson, 2009). Instead of conventional F , t , and p -values, a ratio of the likelihood of the alternative model to the null is reported, usually BF_{10} . For instance, $BF_{10} = 20$ would indicate that the effects model is favored 20 to 1 over the null model. Conversely, if the BF_{10} were 0.10, the null model is favored 10 to 1 over the effects model.

Typical Procedure

The procedure behind BF comparisons requires two steps.

- 1) Similar to the NHST procedure, one must design two models for the data. For our purposes, the first of these models will be the null model, which states that there are no differences between means, and the second model for these analyses is the effects model. The formulas are operationalized as described in the NHST section, however, the assumptions of random and fixed variables are different. Within the Bayesian framework, the observed data X_i are considered to be given, and the parameter μ is considered to be a random variable. Additionally, one must choose the prior distributions that are believed to describe the data. To calculate the posterior distribution, the Bayesian approach considers the prior distribution of our random parameter μ , and updates it using Bayes theorem given the data at hand. Reasonable expectancies of where the data lie should be incorporated in this decision based on previous research into the studied phenomena (Rouder et al., 2012).
- 2) Analyze the data given the selected priors and models. Consider the BF and use the BF_{10} as evidence of how one should update their beliefs about the models.

While the analysis can be quite flexible, the same assumptions that need to be met within the frequentist framework must also be met with the same concerns for attenuated power and inaccurate decisions. Bayesian inference improves upon the traditional frequentist point of view by allowing not only a clear interpretation of the evidence provided by the data, but also the ability to speak in favor of the null hypothesis. It is important to note that while previous work has indicated that p -values and BF largely agree on which hypothesis should be supported, they differ in the strength of that conclusion, especially when p -values were slightly lower than α (i.e., .05 to .01; Wetzels et al., 2011).

As for limitations, Bayesian analyses require the researcher to take an active role in the choice of prior distributions for the phenomenon they are modeling, and this decision can take some effort to fully understand. However, in the meantime, there are packages such as *BayesFactor* that provide the researcher simple default options that can readily lend themselves to many research areas with little fear of being outrageous specifications. Further, unlike NHST, Bayesian analyses do not necessarily control long-run error rates, as the focus is on updating current model beliefs. Another concern that many researchers have is that these analyses are necessarily sensitive to prior choice. Research has shown that the choice of priors has essentially no effect on conclusions when sufficient data has been collected as the priors give way to the weight of the data (Klugkist & Hoijtink, 2007; Rouder et al., 2012), and when reasonable priors are considered, data are only mildly sensitive to these (Haaf & Rouder, 2017). Finally, many believe Bayesian analysis to be too computationally intensive to complete. Many simple programs, packages, and tutorials exist to help ease the transition from frequentist to Bayesian analysis (JASP Team, 2017; Kruschke, 2014; Morey & Rouder, 2015).

Observation Oriented Modeling

History

James Grice argues that our problems as a science go beyond use of NHST and extend into the philosophical ideas underpinning our research. Therefore, he developed a new paradigm called Observation Oriented Modeling (OOM, Grice, 2011, 2014; Grice, Barrett, Schlimgen, & Abramson, 2012). He reasons that by viewing psychology through the lens of philosophical realism, instead of positivism, we should be able to properly and effectively conduct research and analyze data. In contrast to positivism (i.e., which is solely concerned with finding an effect, not with how the effect occurred), philosophical realism holds that the causal structure of nature can be understood through scientific investigation. The goal is then to understand the causal mechanisms that give rise to the patterns observed in a given data set. Switching to this philosophy allows for techniques that match the daily activities of social scientists in their endeavors to unravel the story of how humans operate. OOM pushes the researcher to seek an inference to best explanation (Grice et al., 2017). This causal inference procedure differs from both NHST and Bayes, where a researcher focuses on inferences to population parameters and their various assumptions underlying statistical tests (e.g., random sampling, normality, homogeneity of population treatment differences, etc.).

Generally speaking, the OOM approach can handle any type of data, including ordinal rankings and frequency counts, as all analyses are calculated in the same general fashion (see Valentine and Buchanan (2013) for an example). This simplicity occurs because OOM works on the deep structure of the data which is a binary coding technique similar to dummy and effect coding. Deep structures are matrices of zeros and ones which can be efficiently manipulated to investigate patterns within the data. The most important values from any OOM analysis are the PCC (percent correct classification) values. These values represent the summation of how well each individual's responses matched the stated or expected pattern

or, in the case of causal modeling, how many of the individuals conformed to a given cause.

Complete matches are the proportion of observations that match the researcher-designated pattern on all dimensions. For example, in a three-level Ordinal Pattern Analysis (OPA), a person would be tallied as a “complete match” if the ordinal pattern of their data matched the expected ordinal pattern across all three levels. For example, imagine we have set a pattern that designates Time 1 < Time 2 < Time 3. Person A has values of 3, 4, and 5 at time points 1, 2, and 3, respectively, while person B has values of 4, 5, and 2. Person A matches the pattern completely and is therefore counted in the PCC value. Person B, however, matches only the first part of the pattern (time 1 less than time 2) and is not counted in the PCC value. As the PCC is simply the percentage of individuals in a sample whose responses match the expected ordinal pattern completely, its computation is therefore not based on means or variances, but on the basis of the observations themselves. The PCC value replaces all of the conventional values for effect size used in statistical analyses.

The analysis we focus on here (OPA) does not form any type of linear or nonlinear equation or regression, but simply looks for those individuals who match the expected ordinal pattern (Grice, Craig, & Abramson, 2015). The main point of the analysis, then, is to see how many people fit the expected pattern which is based on a causal theory. If all causes are accounted for in the study and observations have been made with sufficient precision and accuracy, then 100% of the persons should fit the expected pattern; otherwise, a lower PCC value will be expected, and it is up to the researcher to determine how high a PCC must be in order to support an inference to the causal mechanism.

In OOM, traditional p -values are no longer utilized (Grice, 2011). As a secondary form of reference value, a chance value (c -value) is obtained, which is a type of randomization test in which the researcher determines the number of randomized trials for the test (e.g., 1,000 or 5,000 randomized versions of actual observations). This procedure is akin to permutation

tests, where PCCs are computed for the randomized data to form a distribution. The observed PCC is then compared to these values, and the c -value (which is an empirical probability) is determined. If the randomized data sets fit the pattern as well as or better than the actual data does, the c -value will be high (close to 1). Low c -values (close to 0) indicate a pattern of observations that is improbable (i.e., unlikely produced by chance) when compared to randomized versions of the same data. Although low c -values, like low p -values, are desirable, c -values do not adhere to a strict cut-off and should be considered a secondary form of confirmation for the researcher that their results are distinct.

Typical Procedure

OPA is analogous to repeated measures ANOVA and contains two steps.

- 1) Designate the expected ranked pattern: each variable as being higher, lower, or equal to the other variables. For instance, for our analyses we defined the following pattern of individual responses X_i , whereby the first time point should be less than the second time point which should be less than the third time point. This pattern can be operationalized as follows:

$$X_{i_1} < X_{i_2} < X_{i_3}$$

- 2) Analyze the data using OPA. Consider the PCC and c -values in light of the data and use your best judgment as to whether or not the data conform to the expected pattern. This analysis only requires the assumption that the data exists such that a pattern may be designed.

As with all of these methodologies, limitations do exist. This approach is largely concerned with patterns of responses, not with magnitudes of differences, which may be an

integral piece of information to some researchers. Unlike all approaches mentioned before, we do not discuss the probability of some data given our hypothesis here, and instead focus on the observed responses of the individual and how it may or may not behave as expected. Finally, similar to the Bayesian analysis, long-run error rates are not discussed in this methodology.

A Simulation Study

Simulated Data

In this study, we generated 20,000 datasets by manipulating sample size and effect size for a repeated measures design with three levels. A repeated measures design was chosen as it is widely used across many disciplines of psychology. These datasets were created using the *mvtnorm* package in *R* (Genz et al., 2017), and all code for simulations can be found at https://osf.io/u9hf4/?view_only=1caa9092868b4d7aadb9a83a31a979cd. Interested readers can easily adapt the *R* code to incorporate different research designs. Likert data, ranging from 1 to 7, was created by rounding *mvtnorm* estimates to whole numbers and truncating any data points out side of the appropriate range (i.e., values < 1 were rounded to 1, and values > 7 were rounded to 7). We specifically chose Likert-type data as this data type is one of the most common data types utilized by most social scientists. Additionally, we add to the literature as other simulations have chosen to use completely continuous data (i.e., simulated numbers are often precise to 10+ decimals, which is unlikely for traditional sampling). The simulated data did increase in skew with this procedure from approximately no skew (i.e., < 0.01) to approximately 0.40 for the smallest and no effect conditions; however, these values closely resembled a normal distribution with the use of *mvtnorm*. The population means for each level were set to 2.5, 3.0, and 3.5, and pairwise effect sizes (e.g., the comparison between time 1 v. time 2 and time 2 v. time 3) were manipulated by adjusting the standard deviation to create negligible effects ($SD = 3.39$, $d = 0.10$), small

effects ($SD = 3.00$, $d = 0.20$), medium effects ($SD = 0.50$, $d = 0.50$), and large effects ($SD = 0.10$, $d = 0.80$) using Cohen (1992)'s traditional guidelines for d interpretation. The smallest effect size was set such that Likert style data could still be retained with the smallest possible effect size. Sample size was manipulated at 10, 30, 100, 500, and 1,000 data points. All combinations of the five sample sizes and four effect sizes were created, and each dataset was simulated 1,000 times, totaling 20,000 datasets.

The advantage of using *mvtnorm* and set SDs for each group was the ability to approximate the assumptions of normality by randomly generating from a multivariate normal distribution, and homogeneity by setting equal SDs for each group. In a repeated measures design, the assumption of sphericity was met by setting the correlations between levels in *mvtnorm* to zero. By maintaining the lowest level of relationship between levels, we additionally controlled for power and examined situations of significance given the lowest power scenario. During the data simulation, the standard deviation of the difference scores was examined to maintain differences greater than zero, especially for low N simulations.

Analyses Performed

Descriptive Statistics. Means, mean differences between levels, and the confidence intervals for each mean can be found in the complete dataset online, https://osf.io/u9hf4/?view_only=1caa9092868b4d7aadb9a83a31a979cd. For each simulation, we also calculated d values using the standard deviation of the difference score as the denominator (d_z ; Lakens, 2013). The *MOTE* package was used to calculate the non-central confidence interval for each d value as well (Buchanan, Valentine, & Scofield, 2017; Cumming, 2014). This data was mainly used to determine if simulations were meeting expected values overall.

Parametric NHST - Repeated Measures ANOVA. Repeated measures

ANOVA using the *ezANOVA()* function in the *ez* library was utilized with type three sum of squares (Lawrence, 2017). This style of ANOVA is used to compare the same individuals across multiple or all conditions in an experiment. The null hypothesis states that there are no significant differences between population means, and the research hypothesis posits that there are differences between some population means, but does not specify which population means may differ, just that one or more will differ as the alternative. This test uses the F distribution focusing on p values.

To determine where differences may exist, *post hoc* dependent t -tests are normally analyzed in the event of a significant F -ratio. We did not run all pairwise comparisons, instead focusing on the linear trend simulated by comparing level one to two and level two to three. This set of comparisons also controlled the effect size between comparisons, as comparing level one to three would have doubled the effect size. However, we assumed that typical researchers might compare all three pairwise combinations in practice and used a Bonferroni correction across all three possible pairwise combinations to calculate p values for *post hoc* tests. Therefore, while we only discuss the two comparisons, we utilized the more stringent cutoff of the Bonferroni correction as we believe this procedure would be how the majority of researchers would handle the data. Interested readers can find all three comparison values in the complete dataset online. Following traditional usage, a p -value of less than .05 was binned as significant, whereas p -values ranging from .10 to .05 were binned as marginally significant. Any p -values larger than .10 were binned as non-significant. A second set of p -value comparisons was calculated given Benjamin et al. (2018)'s suggestion to change α criterion to less than .005. Any p -value less than .005 was binned as significant, while data ranging from .005 to .10 was marginal or suggestive, and $p > .10$ was non-significant.

Bayesian Analysis: Bayes Factor. We compared a null model with one grand mean for all three levels to an effects model wherein means were allowed to differ using the *BayesFactor* package (Morey & Rouder, 2015). Following Rouder et al. (2012), default priors were placed on the scale of effect sizes (using a g-prior approach). Within the *BayesFactor* package, these were specified by setting the arguments of *rscaleFixed* (prior for standardized fixed effects in the model) and *rscaleRandom* (prior for standardized random effects in the model) to 0.5 and 1.0, respectively. BF were calculated, and follow up *t*-test BFs were computed for the same two comparisons as in the previous models using default priors from the *BayesFactor* package (e.g., Jeffreys prior for population variance, Cauchy prior for standardized effect size). To compare Bayesian results to other statistical methods, we used recommendations from Kass and Raftery (1995) to bin results into weak evidence (BFs < 3), positive evidence (e.g., akin to marginal *p*-values, BFs = 3-20), and strong evidence (BFs > 20). We must stress here that BF interpretation should focus on understanding the odds of model ratios, not necessarily the presence or absence of an effect. However, given that we wanted to compare the conclusions one would reach given this data in a Bayesian paradigm to that of a frequentist paradigm, these bins are used as a convenient comparison to the frequentist procedures using set criteria for interpretation (Morey, 2015). Should any reader become curious how a different set of binning values affect our analyses, all code and data are at their disposal at https://osf.io/u9hf4/?view_only=1caa9092868b4d7aadb9a83a31a979cd, and this manuscript was written with the *papaja* package allowing one to view the code inline with this text (Aust & Barth, 2017).

OOM: Ordinal Pattern Analysis. An *R* script of the Ordinal Pattern Analysis from Grice et al. (2015)'s OOM program was provided from Sauer and Luebke (2016). We set the expected ranked pattern as level one less than level two less than level three. Once this pattern was defined, we then analyzed the data to see if each individual's set of observations matched this expected ordinal pattern. PCC values were generated, and *c*-values were computed by randomizing the data 1,000 times. Solely for purposes of

comparison, we used the following significance coding schema: significant studies had a high PCC value ($.50 < \text{PCC} < 1.00$) and a low c -value ($c < .05$), marginal studies had a high PCC value and a moderate c -value ($.05 < c < .10$), and non-significant studies had low PCC values ($\text{PCC} < .50$), regardless of their c -values. Again, we must stress that this paradigm eschews binning estimates and that our use of bins was a) discussed and decided upon before data analysis, and b) created only for the purposes of comparing this new methodology's possible conclusions to that of a frequentist framework. We welcome interested readers to explore the data more, defining their own bins and viewing the effects, by viewing and editing our code online.

Results

Percent of Estimates

For all simulations, we first binned the estimates into significant, marginal, and non-significant effect categories as described in the Analyses Performed section above. Next, we calculated the percentage of these analyses that would be classified into each of these categories, separated about by statistical analysis, sample size, and effect size. These estimates were binned across both the overall and follow up *post hoc* tests, and the combined data are presented for this analysis. Since all three categories of binning total to 100%, we present only the significant and non-significant results. Significant critical omnibus estimates are presented in Figure 1. All figures discussed in this manuscript may be viewed as interactive graphics on our OSF page through a provided Shiny app. In Figures with sample size on the axes, we log transformed N to allow for visual distinction between sample sizes, as smaller N values were compressed when using the $N = 10$ to 1000 on the axis. Both N and $\log(N)$ can be found in the Shiny app, along with the ability to zoom in to specific ranges of sample size.

For negligible effects at $p < .05$ (solid lines), we found that NSHT analyses showed a predictable Type I error bias, in that they detected significant estimates with extremely small d values as sample size increased. Binned BF values showed a similar pattern, but were more conservative with less percent significant estimates. OOM analyses were the most conservative, essentially never detecting an estimate in the negligible effect simulations. Small effect sizes showed the same pattern for NHST, BF, and OOM results, with the proportion of significant estimates increasing more rapidly and asymptoting at a smaller sample size than negligible effects. At medium effect sizes, NHST analyses nearly always detected significant estimates, while BF and OOM analyses would have been considered significant around 75% of the time. Interestingly, with large effect sizes, OOM analyses mirrored NHST by always detecting estimates, and BF analyses were generally more conservative except at the largest sample size. Figure 1's dashed lines indicate the results if values were binned at $p < .005$, and the differences between these results were very subtle. Lowering α reduced the number of significant estimates at small N values for all four effect sizes, with more pronounced differences at negligible and small effect sizes. However, the graphs converged to the same conclusion that large enough sample sizes could produce significant results at negligible and small effect sizes.

Figure 2 portrays the results for non-significant binned simulations, which were the same for both α criterion. Across all effect sizes, BF and NHST showed similar results, where non-significant estimates were detected at lower sample sizes for negligible and small effect size simulations. At medium and large effect sizes, almost all estimates would have been considered significant, therefore, detection rates for non-significant estimates were around zero. OOM displayed a conservative set of findings, showing nearly 100% non-significant estimates at negligible and small effect sizes (mirroring results from Figure 1). At medium effect sizes, approximately a quarter of estimates were non-significant, illustrating the conservative nature of OOM interpretations.

Figure 3 depicts the relationship between the effect size of time 1 minus time 2 and the corresponding PCC values. These metrics appear to represent different concepts where effect size measures the magnitude of the difference between two data points while PCC disregards magnitude and represents the proportion of the sample following the given ordinal pattern across all three time points. Given these differences, it is interesting how well these two measures converge together. As sample size increases, estimates for both d and PCC become more precise (i.e., smaller range, closer to the simulated effect size). We believe that PCC offers researchers the ability not only to confirm that their effect size is reasonable, but also to better understand the pattern their data are following, especially if an observed effect size contradicts previous literature. For example, let us assume there is previous literature that states that a small positive effect exists, such that responses should increase from time 1 to time 2. Under conditions of a true small effect ($d = -0.20$) and sample size of 30, our graph shows us that it is possible to obtain a positive medium effect size ($d = 0.50$; indicating the time 1 is more extreme than time 2). Upon finding these contradicting results, the researcher could further seek to understand the pattern their data are following by computing the PCC value for the experiment. The PCC value for this example would be above .50, indicating that, in over half of respondents the values for time 1 are less than time 2 (in turn less than time 3, as it measures the entire pattern), even though magnitude of change suggests that time 1 is larger than time 2. This result gives the researcher a richer piece of information, which can help to describe their results in a more nuanced fashion.

Percent Agreement

A goal of this project was to expand the toolbox of options for researchers to determine what evidence supports their hypotheses by examining multiple methodologies. We calculated the percent of time that all analyses agreed across overall and *post hoc* comparison estimates. Figure 4 illustrates the pattern of 100% agreement on effects for critical omnibus

tests only at each sample size and effect size. Figure 5 portrays the results for *post hoc* tests, which only uses NHST and Bayes Factor analyses, as OOM does not have a *post hoc* test (i.e., the test is a pattern analysis that presupposes the expected direction of *post hoc* tests).

When effect sizes were negligible and for small effects, agreement was best across small samples and decreased across sample size, as NHST was overly biased to report significant estimates and OOM and BF were less likely to do so. For medium and large effect sizes, 50-75% agreement was found, usually regardless of sample size. Additionally, we found that for negligible, small, and medium effects, agreement for *post hoc* tests was higher than agreement for overall comparisons. The *post hoc* comparisons for levels 1 to 2 and levels 2 to 3 were less likely to be binned as significant across negligible and small effects, so the agreement levels were higher for these individual comparisons due to non-significant follow up tests. The critical omnibus test was more likely to be significant due to the inclusion of effect of comparisons between level 1 and 3, which were double the effect size. However, these *post hoc* comparisons do not include the conservative significant binning from OOM, which decreased critical omnibus 100% agreement seen in Figure 4. Again, the differences between $p < .05$ and $p < .005$ were minimal. Complete tables of percentages of binning across critical omnibus and *post hoc* tests, along with agreement percentages broken down by bins can be found at https://osf.io/u9hf4/?view_only=1caa9092868b4d7aadb9a83a31a979cd.

Criterion Comparison

As the relationship between BF and p -values is already well documented, we will not discuss them here beyond stating that we found the expected pattern shown in previous work (Rouder et al., 2012), and that individuals who wish to view this comparison, as well as all the other comparisons discussed here should visit our interactive Shiny application at our OSF page. Of interest was the comparison of OOM indices to traditional NHST and Bayesian indices. First, in Figure 6, PCC values are plotted against log BF values and

p-values. The log of BF was taken to include all values on a viewable axis, and all infinity values were windsorized to the next highest point. Increasing sample size is shown by increasing point size and lighter colors. Additionally, since OOM values are a combination of PCC and *c*-values, *c*-values below .05 are shown as Xs instead of dots. Therefore, all values $PCC \geq .50$ that are also denoted as Xs would be considered significant in this example. The provided Shiny application uses color to distinguish sample size differences, as well as includes options to create each combination effect size and criterion individually. Only two graphs are provided here to save space.

In Figure 6, the left hand column portrays the relationship between log BF values and PCC values in negligible and medium effect sizes. With negligible effect sizes, we found large variability in PCC values across a small span of BF values while sample sizes remained low, but as *N* increased, we saw that the range of PCC values narrowed considerably with increasing BF values. Therefore, as sample size increased, the PCC values constricted, while BF values expanded. A similar pattern appeared when viewing the medium sample size graph, as again PCC values became less variable with increased sample size, and BF tended to increase both in variability and in value as the sample size grew. Here, we can see a benefit of PCC, along with *c*-values, as increasing sample size portrayed more precision in PCC, instead of the increased variability found in BF.

It is also important to note that within the negligible effects graph, while many of these PCC values reached high values, that these values did not denote patterns that would necessarily be seen as unique. *c*-values were a secondary measure of evaluation that eliminated a number of these matches from being considered meaningful. A large majority of points with larger sample sizes on the figure included low chance values, however, the PCC values for these simulations were lower than a meaningful percent used for cutoff criterion. This two-step process helped to weed out effects that were negligible, especially at larger sample sizes.

Additionally, we compared p -values and PCC values, which are illustrated on the right hand side of Figure 6. Again, PCC values showed far more variability with small sample sizes, and the p -values associated with these smaller sample sizes were also quite variable. Importantly, even when an effect was negligible, PCC values become less variable with increasing sample size. PCC values also indicated that there was little evidence of the hypothesized pattern by shifting toward zero. p -values decreased in variability at high sample sizes and shifted toward minuscule values, thus, pointing toward rejecting the null hypothesis. With medium effect sizes, both p -values and PCC values were variable at small sample sizes. At larger sample sizes, p -values decreased towards floor effects (i.e., closer to zero), while PCC values simply narrowed in range shifting slight above .50. The benefit of multiple criteria evaluation here was clear, as p -values increasingly indicated significance as sample size increased, PCC values were not effected in this way and thus presented a more stable picture of the presence of an effect. While multiple criteria may not completely reduce the interpretation of false positives in the literature, the relationship between these values illustrated that multiple indices can provided a clearer picture of the evidentiary value available in a study.

Limitations

Within any study a number of limitations exist. The largest limitation of our study is that we chose to focus on a simple three level repeated measures ANOVA design. The benefit to this focus is the simplicity of understanding the relationship between analyses, while also using a well understood NHST procedure. However, is possible that these same relationships may or may not exist in alternative design contexts. Additionally, our choices for classification of significant effects for p -values, BF, PCC, and c -values was based on what we believe a reasonable researcher may designate; however, these classifications may vary in the real world. We provide open access to our simulations and code so that an interested

party can tinker with these choices. We believe the global conclusions would likely be similar across changes, however, the specific percentages and patterns would likely differ. Additionally, as all of our simulations were created within a frequentist framework, this may limit our conclusions regarding they Bayesian methods. Finally, due to the specification of our simulation we did not violate any statistical assumptions. It is possible that the violation of these assumptions may cause changes in the relationships we see here.

Discussion

This manuscript was designed to showcase two alternative paradigms to NHST researchers and to compare the conclusions these alternative methodologies might make in a given data environment to those NHST would make. We believe that the awareness of multiple methodologies might assist in strengthening one's conclusions and improving reproducibility by giving researchers the ability to identify an optimal method given the question at hand. Further, we believe that should a researcher utilize multiple methodologies (e.g., analyzing and reporting both a NHST p -value as well as an OOM PCC value) that these estimates in tandem can help readers to weight these various forms of evidence and arrive at a more robust conclusion. We found that changing the threshold at which p -values are deemed significant had little to no effect on conclusions, especially at large sample sizes, regardless of effect size. This finding is notable as the article by Benjamin et al. (2018) states that an increase in sample size is likely to decrease false positives “by factors greater than two” (p. 10), and work by Pericchi and Pereira (2016) state that an adaptive level of significance would be beneficial in these circumstances, neither of which are supported by our simulations. Our science will not grow by moving the significance line in the sand, as this line has already been shown to have “no ontological basis” (Rosnow & Rosenthal, 1989, p. 1277).

Instead, we need to embrace the multitude of perspectives available to us and to begin to employ these diverse approaches. While NHST can still serve us well when properly

utilized, it is important for researchers to understand that different methods seek to answer different questions, and that we need to ensure that we are using the right method to answer a given question. When evaluating evidence in order to answer these questions we must be wary of looking for significant differences and focus instead on finding meaningful differences. By combining these approaches we may be better able to qualify the strength of our evidence and discuss a more nuanced version of our data. Additionally, while all of these methods have drawbacks, when used in combination these methods can begin to overcome many of these limitations. For instance, given a large sample size, we would expect BF values to be very large and p -values to be very small, both indicating that the null model/hypothesis should not be supported. However, if we also have a PCC value of .30, we may decide that it is possible that this effect is very small and possibly negligible. This multifaceted approach can help to curb our enthusiasm over small or negligible effects that may not be practically meaningful and possibly may not replicate. Regardless if analyses agree or disagree on the presence of an effect, a researcher can investigate the direction and size of the effect, the proportion of data that agrees or disagrees with the direction of the effect, and discuss conclusions accordingly. Each methodology behaves slightly differently in given data environments, which might begin to highlight meaningful differences when discussed together.

Some may contest that all of these analyses are capable of being hacked, like p -values, through researcher degrees of freedom, choice of priors, or pattern choice, among other actions (Simmons et al., 2011). Transparency throughout the research process is key to eliminating these issues, as α changes may only encourage bad research practices with the current incentive structure on publishing. Although we have the capability to share research across the world, research often still occurs behind closed doors. The Open Science Framework grants insight into research processes, allowing researchers to share their methodologies, code, design, and other important components of their projects. In addition to posting materials for projects, pre-registration of hypotheses and methodology will be an important facet in scientific accountability. Further, with increased transparency editors and

other researchers can weigh the evidence presented according to their own beliefs.

Our key suggestion in this project is the redefinition of evidentiary value. The current focus on p -values has shown to be problematic, as many of the studies from the Open Science Collaboration (2015) do not replicate at $p < .05$ or $p < .005$ (Lakens et al., 2018). With the change in transparency mentioned above, publishing research with solid research designs and statistics, regardless of p -values, will allow for a broader range of evidence to become available. Publishing null findings is critical in replication and extension for discovering the limits and settings necessary for phenomena. Registered replications and reports will allow studies to be accepted prior to results being known, thus allowing researchers to focus on experimental design and hypotheses *a priori* instead of p -values *post hoc*. Reports should describe multiple indicators of evidence, such as effect sizes, confidence intervals, power analyses, Bayes Factors, and other descriptive statistics (Finkel, Eastwick, & Reis, 2015; Nosek & Lakens, 2014; van't Veer & Giner-Sorolla, 2016).

A misunderstanding of statistical power still plagues psychological sciences (Bakker, Hartgerink, Wicherts, & van Der Maas, 2016), and the effect of sample size, especially small ones, was shown here by comparing the criterion available in these analyses. Often, individual research labs may not have the means to adequately power a proposed study. Multilab studies and collaboration with other scientists is fundamental to alleviating these issues, while encouraging interdisciplinary science. Collaboration increases our statistical abilities, as every researcher cannot be expected to be proficient in all methods and analyses, but teams of researchers can be assembled to cover a wider range of statistical skills to provide adequate estimates of evidence in their reports. We understand that there may be resistance to the implementation of multiple methodologies as these new methodologies take time and effort to learn. However, through the use of free programs (JASP, *R*, OOM, Shiny) and tutorials (YouTube, Coursera, <http://www.statstools.com>), we believe all researchers are capable of learning these analyses. We believe that through the expansion of our analytic

636 knowledge and application of these new methodologies, we can begin to attenuate some of
637 the strain currently placed on psychological science and to increase the strength of evidence
638 in our discipline.

References

- American Psychological Association. (2010). *Publication manual of the American Psychological Association* (6th ed.). Washington, D.C.: American Psychological Association.
- Aust, F., & Barth, M. (2017). papaja: Create APA manuscripts with R Markdown. Retrieved from <https://github.com/crsh/papaja>
- Bakker, M., Hartgerink, C. H. J., Wicherts, J. M., & van Der Maas, H. L. J. (2016). Researchers' intuitions about power in psychological research. *Psychological Science*, 27(8), 1069–1077. doi:10.1177/0956797616647519
- Bakker, M., van Dijk, A., & Wicherts, J. M. (2012). The rules of the game called psychological science. *Perspectives on Psychological Science*, 7(6), 543–554. doi:10.1177/1745691612459060
- Bellhouse, D. R. (2004). The Reverend Thomas Bayes, FRS: A Biography to celebrate the tercentenary of his birth. *Statistical Science*, 19(1), 3–43. doi:10.1214/088342304000000189
- Benjamin, D. J., Berger, J. O., Johannesson, M., Nosek, B. A., Wagenmakers, E.-J., Berk, R., . . . Johnson, V. E. (2018). Redefine statistical significance. *Nature Human Behaviour*, 2(1), 6–10. doi:10.1038/s41562-017-0189-z
- Berger, J. (2006). The case for objective Bayesian analysis. *Bayesian Analysis*, 1(3), 385–402. doi:10.1214/06-BA115
- Buchanan, E. M., Valentine, K. D., & Scofield, J. E. (2017). MOTE. Retrieved from <https://github.com/doomlab/MOTE>

Cohen, J. (1992). A power primer. *Psychological Bulletin*, 112(1), 155–159.

doi:10.1037/0033-2909.112.1.155

Cumming, G. (2008). Replication and p intervals. *Perspectives on Psychological Science*,

3(4), 286–300. doi:10.1111/j.1745-6924.2008.00079.x

Cumming, G. (2014). The new statistics: Why and how. *Psychological Science*, 25(1), 7–29.

doi:10.1177/0956797613504966

Datta, G., & Ghosh, M. (1996). On the invariance of noninformative priors. *The Annals of*

Statistics, 24(1), 141–159. doi:10.1214/aos/1033066203

De Laplace, P. S. (1774). Mémoire sur les suites récurro-récurrentes et sur leurs usages dans

la théorie des hasards. *Mém. Acad. R. Sci. Paris*, 6(8), 353–371.

Dienes, Z. (2008). *Understanding psychology as a science: an introduction to scientific and*

statistical inference. Palgrave Macmillan.

Dienes, Z. (2014). Using Bayes to get the most out of non-significant results. *Frontiers in*

Psychology, 5(July), 1–17. doi:10.3389/fpsyg.2014.00781

Etz, A., & Vandekerckhove, J. (2016). A Bayesian perspective on the reproducibility project:

Psychology. *PLoS ONE*, 11(2), 1–12. doi:10.1371/journal.pone.0149794

Finkel, E. J., Eastwick, P. W., & Reis, H. T. (2015). Best research practices in psychology:

Illustrating epistemological and pragmatic considerations with the case of relationship

science. *Journal of Personality and Social Psychology*, 108(2), 275–297.

doi:10.1037/pspi0000007

Fisher, R. A. (1932). Inverse probability and the use of Likelihood. *Mathematical*

Proceedings of the Cambridge Philosophical Society, 28(3), 257–261.

doi:10.1017/S0305004100010094

Gelman, A., Carlin, J. B., Stern, H. S., & Rubin, D. R. (2013). *Bayesian data analysis*.
Chapman & Hall/CRC.

Genz, A., Bretz, F., Miwa, T., Mi, X., Leisch, F., Scheipl, F., & Hothorn, T. (2017).
mvtnorm: Multivariate normal and t distributions. Retrieved from
<http://cran.r-project.org/package=mvtnorm>

Gigerenzer, G. (2004). Mindless statistics. *The Journal of Socio-Economics*, 33(5), 587–606.
doi:10.1016/j.socec.2004.09.033

Gigerenzer, G., Krauss, S., & Vitouch, O. (2004). The null ritual: What you always wanted
to know about significance testing but were afraid to ask. In *The sage handbook of
quantitative methodology for the social sciences* (pp. 392–409). Thousand Oaks, CA:
SAGE Publications, Inc. doi:10.4135/9781412986311.n21

Goodman, S. N. (1999). Toward evidence-based medical statistics. 1: The p value fallacy.
Annals of Internal Medicine, 995–1004.
doi:10.7326/0003-4819-130-12-199906150-00008

Grice, J. W. (2011). *Observation Oriented Modeling: Analysis of cause in the behavioral
sciences*. Elsevier/Academic Press.

Grice, J. W. (2014). Observation Oriented Modeling: Preparing students for research in the
21st century. *Comprehensive Psychology*, 3, 05.08.IT.3.3. doi:10.2466/05.08.IT.3.3

Grice, J. W., Barrett, P., Cota, L., Felix, C., Taylor, Z., Garner, S., ... Vest, A. (2017).
Four bad habits of modern psychologists. *Behavioral Sciences*, 7(4), 53.
doi:10.3390/bs7030053

Grice, J. W., Barrett, P. T., Schlimgen, L. A., & Abramson, C. I. (2012). Toward a brighter
future for psychology as an observation oriented science. *Behavioral Sciences*, 2(4),

1–22. doi:10.3390/bs2010001

Grice, J. W., Craig, D. P. A., & Abramson, C. I. (2015). A simple and transparent alternative to repeated measures ANOVA. *SAGE Open*, 5(3), 2158244015604192. doi:10.1177/2158244015604192

Haaf, J., & Rouder, J. N. (2017). *Developing constraint in bayesian mixed models*. doi:10.17605/OSF.IO/KTJNQ

Ioannidis, J. P. A. (2005). Why most published research findings are false. *PLoS Medicine*, 2(8), e124. doi:10.1371/journal.pmed.0020124

JASP Team. (2017). JASP. Retrieved from <https://jasp-stats.org/>

Kass, R. E., & Raftery, A. E. (1995). Bayes Factors. *Journal of the American Statistical Association*, 90(430), 773–795. doi:10.2307/2291091

Klugkist, I., & Hoijsink, H. (2007). The Bayes factor for inequality and about equality constrained models. *Computational Statistics & Data Analysis*, 51(12), 6367–6379. doi:10.1016/j.csda.2007.01.024

Kruschke, J. K. (2014). *Doing Bayesian data analysis: A tutorial with R, JAGS, and Stan* (2nd ed.). Academic Press.

Lakens, D. (2013). Calculating and reporting effect sizes to facilitate cumulative science: A practical primer for t-tests and ANOVAs. *Frontiers in Psychology*, 4. doi:10.3389/fpsyg.2013.00863

Lakens, D. (2017). Equivalence tests. *Social Psychological and Personality Science*, 8(4), 355–362. doi:10.1177/1948550617697177

Lakens, D., Adolfs, F. G., Albers, C. J., Anvari, F., Apps, M. A. J., Argamon, S. E., . . .

Zwaan, R. A. (2018). Justify your alpha. *Nature Human Behaviour*, 2(3), 168–171.
doi:10.1038/s41562-018-0311-x

Lawrence, M. A. (2017). ez: Easy analysis and visualization of factorial experiments.
Retrieved from <http://cran.r-project.org/package=ez>

Lehmann, E. L. (1993). The Fisher, Neyman-Pearson theories of testing hypotheses: One
theory or two? *Journal of the American Statistical Association*, 88(424), 1242–1249.
doi:10.1080/01621459.1993.10476404

Lehmann, E. L. (2011). *Fisher, Neyman, and the creation of classical statistics*. New York,
NY: Springer.

Lindsay, D. S. (2015). Replication in Psychological Science. *Psychological Science*, 26(12),
1827–1832. doi:10.1177/0956797615616374

Maxwell, S. E., & Delaney, H. D. (2004). *Designing experiments and analyzing data: A
model comparison perspective* (2nd ed.). Mahwah, NJ: Lawrence Erlbaum Associates.

Maxwell, S. E., Lau, M. Y., & Howard, G. S. (2015). Is psychology suffering from a
replication crisis? What does “failure to replicate” really mean? *American
Psychologist*, 70(6), 487–498. doi:10.1037/a0039400

Morey, R. D. (2015). On verbal categories for the interpretation of Bayes factors. Retrieved
from [http:
//bayesfactor.blogspot.com/2015/01/on-verbal-categories-for-interpretation.html](http://bayesfactor.blogspot.com/2015/01/on-verbal-categories-for-interpretation.html)

Morey, R. D., & Rouder, J. N. (2015). BayesFactor: Computation of Bayes Factors for
common designs. Retrieved from <https://cran.r-project.org/package=BayesFactor>

Nosek, B. A., & Lakens, D. (2014). Registered reports. *Social Psychology*, 45(3), 137–141.
doi:10.1027/1864-9335/a000192

Nosek, B. A., Spies, J. R., & Motyl, M. (2012). Scientific utopia. *Perspectives on Psychological Science*, 7(6), 615–631. doi:10.1177/1745691612459058

Open Science Collaboration. (2015). Estimating the reproducibility of psychological science. *Science*, 349(6251), aac4716. doi:10.1126/science.aac4716

Pericchi, L., & Pereira, C. (2016). Adaptive significance levels using optimal decision rules: Balancing by weighting the error probabilities. *Brazilian Journal of Probability and Statistics*, 30(1), 70–90. doi:10.1214/14-BJPS257

Press, S. J. (2002). *Subjective and objective Bayesian statistics* (2nd ed.). Hoboken, NJ, USA: John Wiley & Sons, Inc. doi:10.1002/9780470317105

Rosnow, R. L., & Rosenthal, R. (1989). Statistical procedures and the justification of knowledge in psychological science. *American Psychologist*, 44(10), 1276–1284. doi:10.1037/0003-066X.44.10.1276

Rouder, J. N., Morey, R. D., Speckman, P. L., & Province, J. M. (2012). Default Bayes factors for ANOVA designs. *Journal of Mathematical Psychology*, 56(5), 356–374. doi:10.1016/j.jmp.2012.08.001

Rouder, J. N., Speckman, P. L., Sun, D., Morey, R. D., & Iverson, G. (2009). Bayesian t tests for accepting and rejecting the null hypothesis. *Psychonomic Bulletin & Review*, 16(2), 225–237. doi:10.3758/PBR.16.2.225

Sauer, S., & Luebke, K. (2016, January). Observation Oriented Modeling revised from a statistical point of view. doi:10.17605/OSF.IO/3J4XR

Sellke, T., Bayarri, M. J., & Berger, J. O. (2001). Calibration of p values for testing precise null hypotheses. *American Statistician*, 55(1), 62–71. doi:10.1198/000313001300339950

- Simmons, J. P., Nelson, L. D., & Simonsohn, U. (2011). False-positive psychology: Undisclosed flexibility in data collection and analysis allows presenting anything as significant. *Psychological Science*, 22(11), 1359–1366. doi:10.1177/0956797611417632
- Tabachnick, B. G., & Fidell, L. S. (2012). *Using multivariate statistics* (Sixth.). Boston, MA: Pearson.
- Trafimow, D., Amrhein, V., Areshenkoff, C. N., Barrera-Causil, C. J., Beh, E. J., Bilgiç, Y. K., . . . Marmolejo-Ramos, F. (2018). Manipulating the alpha level cannot cure significance testing. *Frontiers in Psychology*, 9. doi:10.3389/fpsyg.2018.00699
- Valentine, K. D., & Buchanan, E. M. (2013). JAM-boree: An application of observation oriented modelling to judgements of associative memory. *Journal of Cognitive Psychology*, 25(4), 400–422. doi:10.1080/20445911.2013.775120
- van Elk, M., Matzke, D., Gronau, Q. F., Guan, M., Vandekerckhove, J., & Wagenmakers, E.-J. (2015). Meta-analyses are no substitute for registered replications: A skeptical perspective on religious priming. *Frontiers in Psychology*, 6, 1365. doi:10.3389/fpsyg.2015.01365
- van't Veer, A. E., & Giner-Sorolla, R. (2016). Pre-registration in social psychology—A discussion and suggested template. *Journal of Experimental Social Psychology*, 67, 2–12. doi:10.1016/j.jesp.2016.03.004
- Wagenmakers, E.-J. (2007). A practical solution to the pervasive problems of p values. *Psychonomic Bulletin & Review*, 14(5), 779–804. doi:10.3758/BF03194105
- Wasserstein, R. L., & Lazar, N. A. (2016). The ASA's statement on p -values: Context, process, and purpose. *The American Statistician*, 70(2), 129–133. doi:10.1080/00031305.2016.1154108

- 798 Wetzels, R., Matzke, D., Lee, M. D., Rouder, J. N., Iverson, G. J., & Wagenmakers, E.-J.
799 (2011). Statistical evidence in experimental psychology. *Perspectives on Psychological*
800 *Science*, 6(3), 291–298. doi:10.1177/1745691611406923

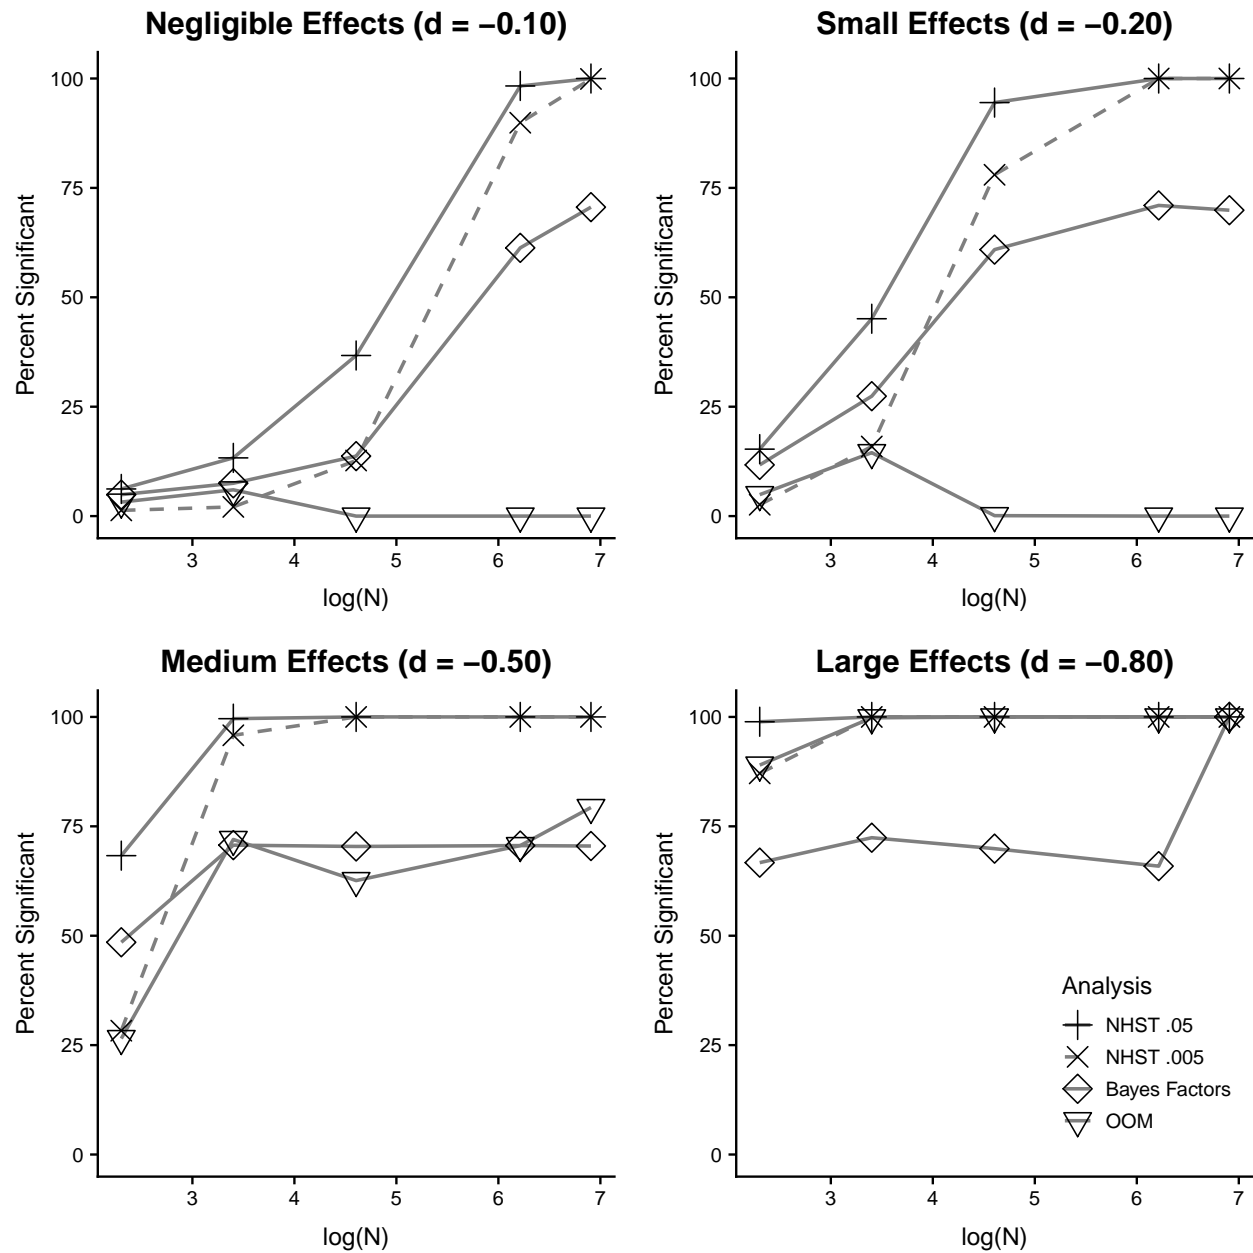


Figure 1. For NHST analyses only, percent of significant estimates at $p < .05$ (solid) and $p < .005$ (dashed) for each analysis given effect size and sample size.




Figure 2. Percent of non-significant effects for each analysis given effect size and sample size.




Figure 3. PCC and c -values plotted against observed effect size (d -values) given effect size and sample size conditions. Xs indicate simulations with c -values $< .05$, which were binned as significant if they were found in conjunction with PCC values over $.50$. Point size and color indicates sample size wherein larger samples are lighter colors. Please note that the key for the entire set of graphics is in the bottom right hand corner.

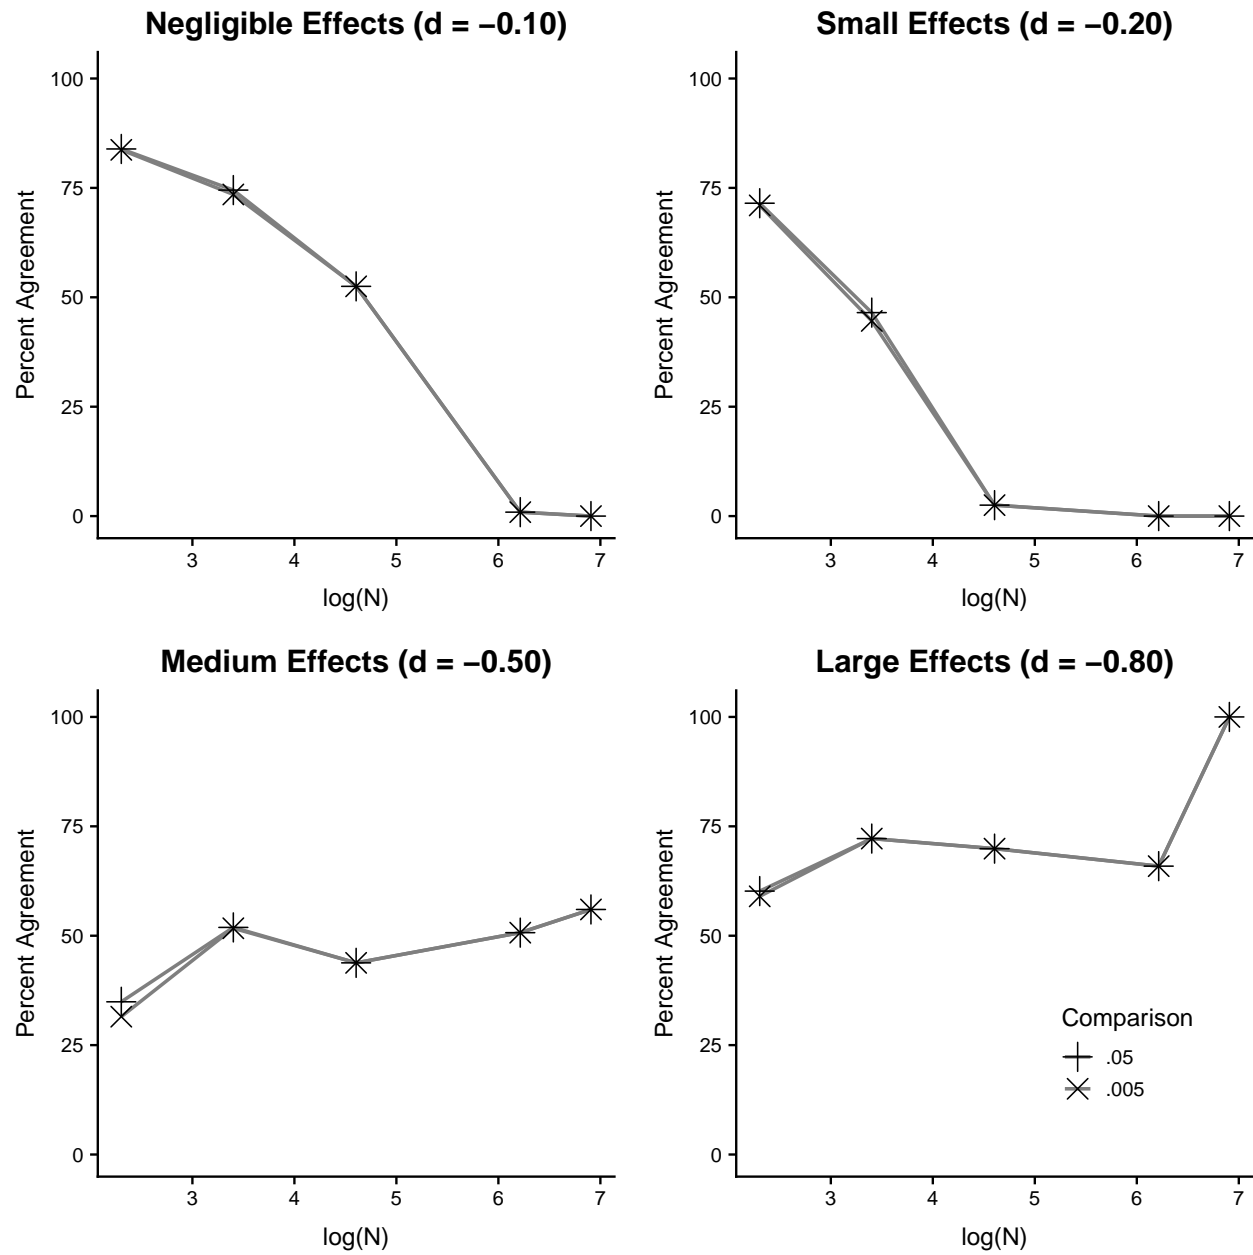


Figure 4. Percent of agreement across all analyses given effect size and sample size for omnibus tests.




Figure 5. Percent of agreement across each analysis given effect size and sample size *posthoc* tests with $p < .05$ (solid) and $p < .005$ (dashed). Note that this graph only compares the NHST and BF conclusions.




Figure 6. PCC and c-values plotted against p and BF values for negligible and medium effect size conditions. Xs indicate simulations with c -values $< .05$, which were binned as significant if they were found in conjunction with PCC values over .50. Point size and color indicates sample size wherein larger samples are lighter colors. Please note that the key for the entire set of graphics is in the bottom right hand corner.