1   Beyond *p*-values: Utilizing Multiple Methods to Evaluate Evidence

2 Kathrene D. Valentine[1], Erin M. Buchanan[2], John E. Scofield[1], & Marshall T. Beauchamp[3]

3   [1] University of Missouri
4   [2] Missouri State University
5   [3] University of Missouri - Kansas City

6   Author Note

Abstract

Null hypothesis significance testing is frequently cited as a threat to the validity and reproducibility of the social sciences. While many individuals suggest we should focus on altering the p-value at which we deem an effect significant, we believe this suggestion is short-sighted. Alternative procedures (i.e., Bayesian analyses and Observation Oriented Modeling; OOM) can be more powerful and meaningful to our discipline. However, these methodologies are less frequently utilized and are rarely discussed in combination with NHST. Herein, we discuss the historical roots, procedures, and assumptions of three methodologies (NHST, Bayesian Model comparison, and OOM), then compare the possible interpretations of three analyses (ANOVA, Bayes Factor, and an Ordinal Pattern Analysis) in various data environments using a simulation study. Our frequentist simulation approach generated 20,000 unique datasets which varied sample size (Ns of 10, 30, 100, 500, 1,000), and effect sizes (ds of 0.10, 0.20, 0.50, 0.80). Through this simulation, we find that changing the threshold at which p-values are considered significant has little to no effect on conclusions. Further, we find that evaluating multiple estimates as evidence of an effect can allow for a more robust and nuanced report of findings. These findings suggest the need to redefine evidentiary value and reporting practices.

Keywords: null hypothesis testing, p-values, Bayes Factors, Observation Oriented Modeling, evidence

Beyond *p*-values: Utilizing Multiple Methods to Evaluate Evidence

Recent events in psychological science have prompted concerns within the discipline regarding research practices and ultimately the validity and reproducibility of published reports (Etz & Vandekerckhove, 2016; Lindsay, 2015; Open Science Collaboration, 2015; Van Elk et al., 2015). One often discussed matter is over-reliance, abuse, and potential hacking of *p*-values produced by frequentist null hypothesis significance testing (NHST), as well as misinterpretations of NHST results (Gigerenzer, 2004; Ioannidis, 2005; Simmons, Nelson, & Simonsohn, 2011). We agree with these concerns and believe that many before us have voiced sound, generally accepted opinions on potential remedies, such as an increased focus on effect sizes (Cumming, 2008; Lakens, 2013; Maxwell, Lau, & Howard, 2015; Nosek, Spies, & Motyl, 2012). However, other suggestions have been met with less enthusiasm, including a recent article by Benjamin et al. (2018) advocating that researchers should begin thinking only of *p*-values less than .005 as "statistically significant", thus changing α levels to control Type I error rates. Additionally, Pericchi and Pereira (2016) promote the use of fluctuating α levels as a function of sample size to assist with these errors. We argue it is not the threshold, or critical *p*-, that needs to be rethought when seeking evidence, but rather if a *p*-value should be utilized at all, and, if so, what that *p*-value can tell you in relation to other indicators. While NHST and *p*-values may have merit, researchers have a wealth of other statistical tools available to them. We believe that improvements may be made to the sciences as a whole when individuals become aware of the tools available to them and how these methods may be used—either alone or in combination—to strengthen understanding and conclusions. These sentiments have been shared by the American Statistical Association who recently held a conference focusing on going beyond NHST, expanding their previous stance *p*-values (Wasserstein & Lazar, 2016).

Therefore, we undertook this project to show researchers how two alternative paradigms compare to NHST in terms of their methodological design, statistical interpretations, and (WHAT WORD SHOULD GO HERE?) robustness. Herein, we will

discuss the following methodologies: NHST, Bayes Factor comparisons, and Observation Oriented Modeling. The three approaches will be compared via this simulated data using a 3 timepoint repeated measures design with a Likert-type scale as the outcome variable. One goal of this study is to introduce social scientists to Observation Oriented Modeling (OOM), as it is a relatively new paradigm that is readily interpretable and, as we will show, useful in these contexts. Additionally, we aim to discuss the conclusions these three methods would make given the same data, and to compare how often these methodologies agree within different data environments (i.e. given varying sample sizes and effect sizes). We hope that by discussing these methodologies in terms of a simple statistical analysis researchers will be able to easily compare and contrast methodologies. For this discussion, it is important to understand their historical background, procedural steps, and limitations, which are outlined below. After this discussion, we describe a simulation study comparing methodologies and α criteria, and end with potential implications for researchers.

Null Hypothesis Significance Testing

History

Many attribute the frequentist NHST procedure to Ronald A. Fisher (Fisher, 1932). However, Fisher's ideas are a far cry from the NHST procedure implemented today. Fisher believed in creating one "null" hypothesis, which he described as a hypothesis to be "nullified", or shown incorrect, not as a zero-difference hypothesis (Lehmann, 2011). He also believed that the use of any omnibus level of significance showed a "lack of statistical thinking" (Gigerenzer, Krauss, & Vitouch, 2004). He instead believed we should report the exact significance value of a test and let others make their own decision about the claims, which is more in line with the typical reporting recommendations provided by the American Psychological Association (American Psychological Association, 2010). Fisher spoke of this work to William Gosset, the man who created the Student's t-test and contributed work on the correlation coefficient (Lehmann, 2011). Gosset in turn discussed the idea of an

alternative hypothesis, a piece not included in Fisher's procedure, with decision theorist

Egon Pearson.

From this discussion, Egon Pearson and Jerzy Neyman created Neyman-Pearson

decision theory. This theory consists of two hypotheses (i.e., null and alternative) and a

binary decision criteria (i.e., significant or not, Lehmann, 1993). However, this procedure

created the possibility of researcher decision errors (Dienes, 2008). A researcher may falsely

reject the null hypothesis (Type I error, α) or falsely fail to reject the null (Type II error, β).

α levels set the binary decision criteria, which are used as the critical p-value for hypothesis

testing (i.e., $p < .05$), and are thus seen as evidence to reject the null hypothesis. β and

power are inherently linked (Power = 1-β), so as the likelihood of finding a true effect

increases beta decreases (Maxwell & Delaney, 2004). Although α values can be chosen to be

quite small, and methods (such as decreasing error variance or using a 1-tailed test as

opposed to a 2-tailed test) can decrease β values as well, a researcher can never know if they

have made the correct decision, or a decision error. Thus, Neyman and Pearson clearly state

that a hypothesis should not be blindly supported based solely on the estimates of one

statistical test, and that replication and reproduction of results are imperative. The recent

work of the Open Science Collaboration (2015) has also highlighted the need for replication

studies and interpretation of results in an appropriate context. Additionally, Neyman and

Pearson emphasized that use of set αs and βs is illogical and sought instead for researchers

to adjust their analysis to the needs of the particular task at hand (Gigerenzer, 2004).

Typical NHST Procedure

Neither Fisher's hypothesis testing, nor Neyman-Pearson decision theory quite match

the NSHT procedure as it is taught and applied today. Psychologists have largely adopted

an amalgamation of the two approaches. Here, we attempt to outline what we believe is the

most appropriate way to carry out the traditional NHST procedure in the context of a

repeated measures ANOVA with three levels, although we note that this set of steps is not

115 necessarily how researchers carry out the procedure in practice:

116   1) Create two hypotheses, one to be "nullified" and one "alternative" hypothesis. Within

117   this repeated measures framework, most researchers would define a null hypothesis

118   (H_0) that indicates of all three time point population means are equal. The alternative

119   hypothesis (H_A) would then be that not all of the population means are equal. These

120   can be operationalized in our example data as follows (note that for H_A we use a

121   common short hand to denote the model):

$$H_0 : \mu_1 = \mu_2 = \mu_3$$

$$H_A : \mu_1 \neq \mu_2 \neq \mu_3$$

122 2) Select an α level that is appropriate given the context of your research, your analysis plan,

123 and your research question, and do not blindly adopt an omnibus critical p-value (Lakens et

124 al., 2018).

125   3) Compute your given analysis and identify the corresponding p-value. If your p-value is

126   less than the chosen α, reject the null hypothesis and state that there appear to be

127   differences between some of your population means; however, if your p-value is greater

128   than or equal to the value selected, do not reject the null hypothesis, and state that a

129   difference between the population means could not be supported.

130   While the NHST procedure itself gives us testable models, the specific analysis used to

131 test these models here, the repeated measures ANOVA with 3 levels, requires some

132 additional assumptions that must be met before an analysis is begun (Tabachnick & Fidell,

133 2012). Data need to have no outlying or influential observations. Data must have a normal

134 sampling distribution, be linearly related, and have independent errors. Depending on the

135 statistical test, data must also be checked for equal variances, sphericity, and additivity.

136 These assumptions can be checked and, if necessary, corrected for; however, violations of

these assumptions can lead to inaccurate decisions and attenuated power. Further, with many analysis programs, data is required to have no missing values.

While this approach is widely used, there are many limitations associated with it. First, this method can be sensitive to violations of the stated assumptions, and especially, if the sample size is not large enough to create a normal sampling distribution (Tabachnick & Fidell, 2012). Even if assumptions are met, or nonparametric tests are implemented (e.g., for situations where a normal distribution assumption cannot be met), this methodology does not allow a researcher to state anything about the absence of an effect (i.e., no true differences). Through traditional NHST, one can only discuss evidence regarding the alternative hypothesis; one can never support the null hypothesis through this procedure. Given the recent findings regarding reproducibility, showing support for the absence of an effect can be even more crucial than showing support for the presence of an effect (Bakker, Van Dijk, & Wicherts, 2012; Lakens, 2017).

Bayes Factors

History

Thomas Bayes was a statistician and Presbyterian minister whose works are still influential today (Bellhouse, 2004). Bayes' theorem solved the inverse probability problem, namely that through the frequentist approach, one can only know the probability of data existing given a hypothesis being true, never the probability of a hypothesis being true given that the data exist (Dienes, 2008). Bayes' theorem allows one to calculate the probability of a hypothesis given some data (posterior belief) by using how probable one believes the hypothesis to be before data was collected (prior belief) and how probable one believes the data to be given one's hypothesis (likelihood). Thus, with his theorem, researchers are able to update (through the use of the likelihood) our initial beliefs (our prior) given some data (Gelman, Carlin, Stern, & Rubin, 2013). Pierre-Simon Laplace pioneered Bayesianism and advocated for a broader interpretation of this theorem (De Laplace, 1774). The use of

Bayesian statistics has been suggested as an NHST alternative (Dienes, 2014; Wagenmakers, 2007), but this approach has largely been undervalued in favor of frequentist methods as, until recently, Bayesian analysis required considerable computational effort. However, today we possess the technology necessary to efficiently conduct Bayesian analyses. While open source software, such as *R* and JASP, require minimal learning to be able to effectively operate (Morey & Rouder, 2015), researchers will need to invest more effort to understand the focus and interpretation of Bayes Factor (BF) comparisons as they differ from traditional NHST.

The Bayesian framework can be viewed as a continuum, with objective Bayesian analyses on one end, and subjective Bayesian analyses on the other (Press, 2002). While this topic could lend itself to its own manuscript, here we will simply summarize the two endpoints, and discuss where our analysis may be perceived to fall on the line. Objective Bayesian analysis is closest to frequentist theory, as the aim is to minimize the influence of priors through the use of non-informative priors (such as Jefferys priors that are designed to be invariant under reparameterization Datta & Ghosh, 1996); thus, the data is allowed to maximally effect the posterior distribution. Further, objective Bayesian methods are influenced by the same quality criteria that frequentist methods used, including Type I error rate and power (Sellke, Bayarri, & Berger, 2001). On the other end, subjective Bayes analyses include rigorously informed priors so that current knowledge can play a large role in the posterior. Our current analysis splits these two; we do not utilize completely uniformed (objective) priors, as we can adjust for basic knowledge of the constraints of our data type. Given the usual lack of information about underlying distributions, a wider band of inclusion was used for prior information. The *BayesFactor* package (Morey & Rouder, 2015) assists greatly in the choice of prior and is especially user-friendly for applied researchers, as it makes use of recommended default priors that have been chosen to be safe to assume under a broad range of data and topics (Rouder, Morey, Speckman, & Province, 2012; Rouder, Speckman, Sun, Morey, & Iverson, 2009). Instead of conventional *F*, *t*, and *p*-values, a ratio

of the likelihood of the alternative model to the null is report, usually BF_{10}. For instance, $BF_{10} = 20$ would indicate that the effects model is favored 20 to 1 over the null model. Conversely, if the BF_{10} were 0.10, the null model is favored 10 to 1 over the effects model.

Typical Procedure

The procedure behind BF comparisons requires two steps.

1) One must design two models for the data. For our purposes, the first of these models will be the null model, which states that there are no differences between means (μ; i.e. all of our observed values X_i, regardless of which time point they were assessed at X_{ij}, arise from a normal distribution N with some mean μ and variance σ^2). The second model for these analyses is the effects model, which states that each mean (μ) is allowed to be different from the grand mean by some amount (α; as we now have observations being drawn from three potential normal distributions, all of which may have a different mean value, but the same variance). These can be operationalized as follows: CHANGE EFFECT (ALPHA) TO A DIFFERENT LETTER SO AS NOT TO CONFUSE WITH CUTOFF?

$$H_0 : X_{ij} \sim N(\mu, \sigma^2)$$
$$H_A : X_{ij} \sim N(\mu + \alpha_i, \sigma^2)$$

In designing these models, one must choose the prior distributions that are believed to describe the data. Reasonable expectancies of where the data lie should be incorporated in this decision based on previous research into the studied phenomena (Rouder et al., 2012).

2) Analyze the data given the selected priors and models. Consider the BF and use the BF_{10} as evidence of how one should update their beliefs about the models.

Based on the flexibility of the analysis, the only assumption that needs to be made is that data exists such that two competing, plausible models with different constraints may be specified.

Bayesian inference improves upon the traditional frequentist point of view by allowing not only a clear interpretation of the evidence provided by the data, but also the ability to speak in favor of the null hypothesis. It is important to note that while previous work has indicated that p-values and BF largely agree on which hypothesis should be supported, they differ in the strength of that conclusion, especially when p-values were slightly lower than α (i.e., .05 to .01; Wetzels et al., 2011). However, some limitations do arise in this paradigm. Bayesian analyses require the researcher to take an active role in the choice of prior distributions for the phenomenon they are modeling, and this decision can take some effort to fully understand; however, in the meantime, there are packages such as *BayesFactor* that provide the researcher simple default options that can readily lend themselves to many research areas with little fear of being outrageous specifications. Further, unlike NHST, Bayesian analyses do not necessarily control long-run error rates, as the focus is on updating current model beliefs. Another concern that many researchers have is that these analyses are necessarily sensitive to prior choice. However, research has shown that the choice of priors has essentially no effect on conclusions when sufficient data has been collected as the priors give way to the weight of the data (Klugkist & Hoijtink, 2007; Rouder et al., 2012) and when reasonable priors are considered, data are only mildly sensitive to these (Haaf & Rouder, 2017). Finally, many believe Bayesian analysis to be too computationally intensive to complete. However, many simple programs, packages, and tutorials exist to help ease the transition from frequentist to Bayesian analysis (JASP Team, 2017; Kruschke, 2014; Morey & Rouder, 2015).

Observation Oriented Modeling

History

James Grice argues that our problems as a science go beyond use of NHST and extend into the philosophical ideas underpinning our research. Therefore, he developed a new paradigm called Observation Oriented Modeling (OOM, Grice, 2011, 2014; Grice, Barrett,

Schlimgen, & Abramson, 2012). He reasons that by viewing psychology through the lens of philosophical realism, instead of positivism, we should be able to properly and effectively conduct research and analyze data. In contrast to positivism (i.e., which is solely concerned with finding an effect, not with how the effect occurred), philosophical realism holds that the causal structure of nature can be understood through scientific investigation. The goal is then to understand the causal mechanisms that give rise to the patterns observed in a given set of observations, which in here would refer to data. Switching to this philosophy allows for techniques that match the daily activities of social scientists in their endeavors to unravel the story of how humans operate. Using OOM, a researcher does not focus on population parameters and the various assumptions underlying statistical tests (e.g., random sampling, normality, homogeneity of population treatment differences, etc.).

Generally speaking, this approach can handle any type of data, including ordinal rankings and frequency counts, as all analyses are calculated in the same general fashion (see Valentine & Buchanan, 2013 for an example). This simplicity occurs because OOM works on the deep structure of the data. Through observational definition, the program separates these units into binary code. Deep structures can be arranged to form a matrix, which can then be manipulated via matrix algebra, binary Procrustes rotation, and other operations to investigate the data. The most important values from any OOM analysis are the PCC (percent correct classification) values. These values represent the summation of how well an individual's responses matched the stated or expected pattern or, in the case of causal modeling, how many of the individual's conformed to a given cause. Complete matches are the proportion of observations that match the researcher-designated pattern on all dimensions. For example, in a three-level Ordinal Pattern Analysis (OPA), a person would be tallied as a "complete match" if the ordinal pattern of his/her data matched the expected ordinal pattern across all three levels. Imagine we have set a pattern that designates Time 1 < Time 2 < Time 3. For example, imagime we have data for two hypothetical individuals. Person A has values of 3, 4, and 5 at timepoints 1, 2, and 3, respectively, while person B has

values of 4, 5, and 2. We can see that Person A matched the pattern completely, and

therefore would be counted in the PCC value. However, while person B matched the first

part of our pattern (time 1 less than time 2), they did not match on the third point of our

pattern (time 2 less than time 3); thus, they would not be counted in the PCC value. As the

PCC is simply the percentage of individuals in a sample whose responses match the expected

ordinal pattern perfectly, its computation is therefore not based on means or variances, but

on the basis of the observations themselves. The PCC value replaces all of the conventional

values for effect size used in statistical analyses.

The analysis we focus on here (OPA) does not form any type of linear or nonlinear

equation or regression, but simply looks for those individuals who match the expected

ordinal pattern (Grice, Craig, & Abramson, 2015). The main point of the analysis, then, is

to see how many people fit the expected pattern which is based on a causal theory. If all

causes are accounted for in the study and observations have been made with sufficient

precision and accuracy, then 100% of the persons should fit the expected pattern; otherwise,

a lower PCC value will be expected and it is up to the researcher to determine how high a

PCC must be in order to support an inference to the causal mechanism.

In OOM, traditional p-values are no longer utilized (Grice, 2011). As a secondary form

of reference value, a chance value (c-value) is obtained, which is a type of randomization test

in which the researcher determines the number of randomized trials for the test (e.g. 1,000 or

5,000 randomized versions of actual observations). This procedure is akin to permutation

tests, where PCCs are computed for the randomized data to form a distribution. The

observed PCC is then compared to these values, and the c-value (which is an empirical

probability) is determined. If the randomized data sets fit the pattern as well as or better

than the actual data does, the c-value will be high (close to 1). Low c-values (close to 0)

indicate a pattern of observations that is improbable (i.e., unlikely produced by chance)

when compared to randomized versions of the same data. Although low c-values, like low

p-values, are desirable, c-values do not adhere to a strict cut-off and should be considered a

secondary form of confirmation for the researcher that their results are distinct.

Typical Procedure

The OPA is analogous to repeated measures ANOVA and contains two steps.

1) Designate the expected ranked pattern: each variable as being higher, lower, or equal to the other variables. For instance, for our analyses we defined the following pattern of individual responses X_i, whereby the first time point should be less than the second time point which should be less than the third time point. This pattern can be operationalized as follows:

$$X_{i_1} < X_{i_2} < X_{i_3}$$

2) Analyze the data using the OPA. Consider the PCC (the percentage of individuals matching the ordinal hypothesis) and c-values in light of the data and use your best judgment as to whether or not the data conform to the expected pattern. This analysis only requires the assumption that the data exists such that a pattern may be designed.

As with all of these methodologies, limitations do exist. This approach is largely concerned with patterns of responses, not with magnitudes of differences, which may be an integral piece of information to some researchers. Unlike all approaches mentioned before, we do not discuss the probability of some data given our hypothesis here, and instead focus on the observed responses of the individual and how it may or may not behave as expected. Finally, similar to the Bayesian analysis, long-run error rates are not discussed in this methodology.

A Simulation Study

Simulated Data

In this study, we generated 20,000 datasets by manipulating sample size and effect size for a repeated measures design with three levels. A repeated measures design was chosen as it is widely used across many disciplines of psychology. These datasets were created using the *mvtnorm* package in *R* (Genz et al., 2017), and all code for simulations can be found at https://osf.io/u9hf4/. Interested readers can easily adapt the *R* code to incorporate different research designs. Likert data, ranging from 1 to 7, was created by rounding *mvtnorm* estimates to whole numbers and truncating any data points out side of the appropriate range (i.e. values < 1 were rounded to 1, and values > 7 were rounded to 7). We specifically chose Likert-type data as this data type is one of the most common data types utilized by most social scientists. Additionally, we add to the literature as other simulations have chosen to use completely continuous data (i.e., simulated numbers are often precise to 10+ decimals, which is unlikely for traditional sampling). The simulated data did increase in skew with this procedure from approximately no skew (i.e., <0.01) to approximately 0.40 for the smallest and no effect conditions; however, these values closely resembled a normal distribution with the use of *mvtnorm*. The population means for each level were set to 2.5, 3.0, and 3.5, and pairwise effect sizes (e.g., the comparison between time 1 v. time 2 and time 2 v. time 3) were manipulated by adjusting the standard deviation to create negligible effects ($SD = 3.39$, $d = 0.10$), small effects ($SD = 3.00$, $d = 0.20$), medium effects ($SD = 0.50$, $d = 0.50$), and large effects ($SD = 0.10$, $d = 0.80$) using Cohen (1992)'s traditional guidelines for d interpretation. The smallest effect size was set such that Likert style data could still be retained with the smallest possible effect size. Sample size was manipulated at 10, 30, 100, 500, and 1,000 data points. All combinations of the five sample sizes and four effect sizes were created, and each dataset was simulated 1,000 times, totaling 20,000 datasets.

The advantage of using *mvtnorm* and set *SDs* for each group was the ability to approximate the assumptions of normality by randomly generating from a multivariate

normal distribution, and homogeneity by setting equal *SDs* for each group. In a repeated

measures design, the assumption of sphericity was met by setting the correlations between

levels in *mvtnorm* to zero. By maintaining the lowest level of relationship between levels, we

additionally controlled for power and examined situations of significance given the lowest

power scenario. During the data simulation, the standard deviation of the difference scores

was examined to maintain differences greater than zero, especially for low N simulations.

Analyses Performed

   Descriptive Statistics.   Means, mean differences between levels, and the confidence

intervals for each mean can be found in the complete dataset online, https://osf.io/u9hf4/.

For each simulation, we also calculated d values using the standard deviation of the

difference score as the denominator (d_z, Lakens, 2013). The *MOTE* library was used to

calculate the non-central confidence interval for each d value as well (Buchanan, Valentine, &

Scofield, 2017; Cumming, 2014). This data was mainly used to determine if simulations were

meeting expected values overall.

   Parametric NHST - Repeated Measures ANOVA.   Repeated measures

ANOVA using the *ezANOVA()* function in the *ez* library was utilized with type three sum of

squares (Lawrence, 2017). This style of ANOVA is used to compare the same individuals

across multiple or all conditions in an experiment. The null hypothesis states that there are

no significant differences between population means, and the research hypothesis posts that

there are differences between some population means, but does not specify which population

means may differ, just that one or more will differ as the alternative. This test uses the F

distribution focusing on p values.

   To determine where differences may exist, *post hoc* dependent t-tests are normally

analyzed in the event of a significant F-ratio. We did not run all pairwise comparisons,

instead focusing on the linear trend simulated by comparing level one to two and level two to

three. This set of comparisons also controlled the effect size between comparisons, as

comparing level one to three would have doubled the effect size. However, we assumed that typical researchers might compare all three pairwise combinations in practice and used a Bonferroni correction across all three possible pairwise combinations to calculate p values for *post hoc* tests. Therefore, while we only discuss the two comparisons, we utilized the more stringent cutoff of the Bonferroni correction as we believe this procedure would be how the majority of researchers would handle the data. Interested readers can find all three comparison values in the complete dataset online. Following traditional usage, a p-value of less than .05 was binned as significant, whereas p-values ranging from .10 to .05 were binned as marginally significant. Any p-values larger than .10 were binned as non-significant. A second set of p-value comparisons was calculated given Benjamin et al. (2018)'s suggestion to change α criterion to less than .005. Any p-value less than .005 was binned as significant, while data ranging from .005 to .10 was marginal or suggestive, and $p > .10$ was non-significant.

Bayesian Analysis: Bayes Factor. We compared a null model with one grand mean for all three levels to an effects model wherein means were allowed to differ using the *BayesFactor* package (Morey & Rouder, 2015). The default in this package is a Jeffreys prior with a fixed rscale (0.5) and random rscale (1.0). BF were calculated, and follow up t-test BFs were computed for the same two comparisons as in the previous models using default priors from the *BayesFactor* package (e.g., Jeffreys prior for population variance, Cauchy prior for standardized effect size). To compare Bayesian results to other statistical methods, we used recommendations from (**???**) to bin results into weak evidence (BFs < 3), positive evidence (e.g., akin to marginal p-values, BFs = 3-20), and strong evidence (BFs > 20). We must stress here that BF interpretation should focus on understanding the odds of model ratios, not necessarily the presence or absence of an effect. However, given that we wanted to compare the conclusions one would reach given this data in a Bayesian paradigm to that of a frequentist paradigm, these bins are used as a convenient comparison to the freuqntist procedures using set criteria for interpretation (Morey, 2015). Should any reader become

curious how a different set of binning values affect our analyses, all code and data are at their disposal at https://osf.io/u9hf4/.

OOM: Ordinal Pattern Analysis. An *R* script of the Ordinal Pattern Analysis from Grice et al. (2015)'s OOM program was provided from Sauer and Luebke (2016). We set the expected ranked pattern as level one less than level two less than level three. Once this pattern was defined, we then analyzed the data to see if each individual's set of observations matched this expected ordinal pattern. PCC values were generated, and *c*-values were computed by randomizing the data 1,000 times. Solely for purposes of comparison, we used the following significance coding schema: significant studies had a high PCC value ($.50 < \text{PCC} < 1.00$) and a low *c*-value ($c < .05$), marginal studies had a high PCC value and a moderate *c*-value ($.05 < c < .10$), and non-significant studies had low PCC values ($\text{PCC} < .50$), regardless of their *c*-values. Again, we must stress that this paradigm eschews binning estimates and that our use of bins was a) discussed and decided upon before data analysis, and b) created only for the purposes of comparing this new methodologies possible conclusions to that of a frequentist framework. We welcome interested readers to explore the data more, defining their own bins and viewing the affects, by viewing and editing our code at https://osf.io/u9hf4/.

Results

Percent of Estimates

For all simulations, we first binned the estimates into significant, marginal, and non-significant effect categories as described in the Analyses Performed section above. Next, we calculated the percentage of these analyses that would be classified into each of these categories, separated about by statistical analysis, sample size, and effect size. These estimates were binned across both the overall and follow up *post hoc* tests, and the combined data is presented for this analysis. Since all three categories of binning total to 100%, we present only the significant and non-significant results. All analyses and findings can be

found online at https://osf.io/u9hf4/. Significant critical omnibus estimates are presented in Figure 1. All figures discussed in this manuscript may be viewed as interactive graphics on our OSF page through a provided Shiny app. In Figures with sample size on the axes, we log transformed N to allow for visual distinction between sample sizes, as smaller N values were compressed when using the $N = 10$ to 1000 on the axis. Both N and $\log(N)$ can be found in the Shiny app, along with the ability to zoom in to specific ranges of sample size.

For negligible effects at $p < .05$ (solid lines), we found that NSHT analyses showed a predictable Type I error bias, in that they detected significant estimates with extremely small d values as sample size increased. Binned BF values showed a similar pattern, but were more conservative with less percent significant estimates. OOM analyses were the most conservative, essentially never detecting an estimate in the negligible effect simulations. Small effect sizes showed the same pattern for NHST, BF, and OOM results, with the proportion of significant estimates increasing more rapidly and asymptoting at a smaller sample size than negligible effects. At medium effect sizes, NHST analyses nearly always detected significant estimates, while BF and OOM analyses would have been considered significant around 75% of the time. Interestingly, with large effect sizes, OOM analyses mirrored NHST by always detecting estimates, and BF analyses were generally more conservative except at the largest sample size. Figure 1's dashed lines indicate the results if values were binned at $p < .005$, and the differences between these results were very subtle. Lowering α reduced the number of significant estimates at small N values for all four effect sizes, with more pronounced differences at negligible and small effect sizes. However, the graphs converged to the same conclusion that large enough sample sizes could produce significant results at negligible and small effect sizes.

Figure 2 portrays the results for non-significant binned simulations, which were the same for both α criterion. Across all effect sizes, BF and NHST showed similar results, where non-significant estimates were detected at lower sample sizes for negligible and small effect size simulations. At medium and large effect sizes, almost all estimates would have

been considered significant, therefore, detection rates for non-significant estimates were around zero. OOM displayed a conservative set of findings, showing nearly 100% non-significant estimates at negligible and small effect sizes (mirroring results from Figure 1). At medium effect sizes, approximately a quarter of estimates were non-significant, illustrating the conservative nature of OOM interpretations.

Percent Agreement

A goal of this project was to expand the toolbox of options for researchers to determine what evidence supports their hypotheses by examining multiple methodologies. We calculated the percent of time that all analyses agreed across overall and *post hoc* comparison estimates. Figure 4 illustrates the pattern of 100% agreement on effects for critical omnibus tests only at each sample size and effect size. Figure 5 portrays the results for *post hoc* tests, which only uses NHST and Bayes Factor analyses, as OOM does not have a *post hoc* test (i.e., the test is a pattern analysis that presupposes the expected direction of *post hoc* tests).

When effect sizes were negligible and for small effects, agreement was best across small samples and decreased across sample size, as NHST was overly biased to report significant estimates and OOM and BF were less likely to do so. For medium and large effect sizes, 50-75% agreement was found, usually regardless of sample size. Additionally, we found that for negligible, small, and medium effects, agreement for *post hoc* tests was higher than agreement for overall comparisons. The *post hoc* comparisons for levels 1 to 2 and levels 2 to 3 were less likely to be binned as significant across negligible and small effects, so the agreement levels were higher for these individual comparisons due to non-significant follow up tests. The critical omnibus test was more likely to be significant due to the inclusion of effect of comparisons between level 1 and 3, which were double the effect size. However, these *post hoc* comparisons do not include the conservative significant binning from OOM, which decreased critical omnibus 100% agreement seen in Figure 4. Again, the differences between $p < .05$ and $p < .005$ were minimal. Complete tables of percentages of binning

across critical omnibus and *post hoc* tests, along with agreement percentages broken down by bins can be found at https://osf.io/u9hf4/.

Criterion Comparison

As the relationship between BF and *p*-values is already well documented, we will not discuss them here beyond stating that we found the expected pattern shown in previous work (Rouder et al., 2012), and that individuals who wish to view this comparison, as well as all the other comparisons discussed here should visit our interactive Shiny application at our OSF page. Of interest was the comparison of OOM indices to traditional NHST and Bayesian indices. First, in Figure 6, PCC values are plotted against log BF values and *p*-values. The log of BF was taken to include all values on a viewable axis, and all infinity values were windsorized to the next highest point. Increasing sample size is shown by increasing point size and lighter colors. Additionally, since OOM values are a combination of PCC and *c*-values, *c*-values below .05 are shown as Xs instead of dots. Therefore, all values PCC >= .50 that are also denoted as Xs would be considered significant in this example. The provided Shiny application uses color to distinguish sample size differences, as well as includes options to create each combination effect size and criterion individually. Only two graphs are provided here to save space.

In Figure 6, the left hand column portrays the relationship between log BF values and PCC values in negligible and medium effect sizes. With negligible effect sizes, we found large variability in PCC values across a small span of BF values while sample sizes remained low, but as *N* increased, we saw that the range of PCC values narrowed considerably with increasing BF values. Therefore, as sample size increased, the PCC values constricted, while BF values expanded. A similar pattern appeared when viewing the medium sample size graph, as again PCC values became less variable with increased sample size, and BF tended to increase both in variability and in value as the sample size grew. Here, we can see a benefit of PCC, along with *c*-values, as increasing sample size portrayed more precision in

PCC, instead of the increased variability found in BF.

It is also important to note that within the negligible effects graph, while many of these PCC values reached high values, that these values did not denote patterns that would necessarily be seen as unique. *c*-values were a secondary measure of evaluation that eliminated a number of these matches from being considered meaningful. A large majority of points with larger sample sizes on the figure included low chance values, however, the PCC values for these simulations were lower than a meaningful percent used for cutoff criterion. This two-step process helped to weed out effects that were negligible, especially at larger sample sizes.

Additionally, we compared *p*-values and PCC values, which are illustrated on the right hand side of Figure 6. Again, PCC values showed far more variability with small sample sizes, and the *p*-values associated with these smaller sample sizes were also quite variable. Importantly, even when an effect was negligible, PCC values become less variable with increasing sample size. PCC values also indicated that there was little evidence of the hypothesized pattern by shifting toward zero. *p*-values decreased in variability at high sample sizes and shifted toward minuscule values, thus, pointing toward rejecting the null hypothesis. With medium effect sizes, both *p*-values and PCC values were variable at small sample sizes. At larger sample sizes, *p*-values decreased towards floor effects (i.e. closer to zero), while PCC values simply narrowed in range shifting slight above .50. The benefit of multiple criteria evaluation here was clear, as *p*-values increasingly indicated significance as sample size increased, PCC values were not effected in this way and thus presented a more stable picture of the presence of an effect. While multiple criteria may not completely reduce the interpretation of false positives in the literature, the relationship between these values illustrated that multiple indices can provided a clearer picture of the evidentiary value available in a study. CLARIFY HERE THAT WE'RE NOT SAYING TO USE ALL THESE AT THE SAME TIME, BUT THAT RESEARCHERS SHOULD CONSIDER THAT THEY HAVE MULTIPLE OPTIONS AND KNOW THE BENEFITS, DRAWBACKS, AND HOW

525 EACH COMPARE.

Limitations

527   Within any study a number of limitations exist. The largest limitation of our study is

528 that we chose to focus on a simple three level repeated measures ANOVA design. The benefit

529 to this focus is the simplicity of understanding the relationship between analyses, while also

530 using a well understood NHST procedure. However, is possible that these same relationships

531 (CLARIFY THAT THE RELATIONSHIPS WE ARE TALKING ABOUT ARE HOW

532 THESE 3 METHODS ARE SIMILAR/DIFFERENT/AGREE/DISAGREE) may or may not

533 exist in alternative design contexts. Additionally, our choices for classification of significant

534 effects for p-values, BF, PCC, and c-values was based on what we believe a reasonable

535 researcher may designate; however, these classifications may vary in the real world. We

536 provide open access to our simulations and code so that an interested party can tinker with

537 these choices. We believe the global conclusions would likely be similar across changes,

538 however, the specific percentages and patterns would likely differ. Finally, due to the

539 specification of our simulation we did not violate any statistical assumptions. It is possible

540 that the violation of these assumptions may cause changes in the relationships we see here.

Discussion

542   This manuscript was designed to showcase two alternative paradigms to NHST

543 researchers and to compare the conclusions these alternative methodologies might make in a

544 given data environment to those NHST would make. We believe that the awareness of

545 multiple methodologies might assist in strengthening our conclusions and improving

546 reproducibility by giving researchers the ability to identify an optimal method given the

547 question at hand. Further, we believe that should a researcher utilize multiple methodologies

548 (e.g., analyzing and reporting both a NHST p-value as well as an OOM PCC value) that

549 these estimates in tandem can help readers to weight these various forms of evidence and

550 arrive at a more robust conclusion. We found that changing the threshold at which p-values

are deemed significant had little to no effect on conclusions, especially at large sample sizes, regardless of effect size. This finding is notable as the article by Benjamin et al. (2018) states that an increase in sample size is likely to decrease false positives "by factors greater than two" (p. 10), and work by Pericchi and Pereira (2016) state that an adaptive level of significance would be beneficial in these circumstances—neither of which are not supported by our simulations. Our science will not grow by moving the significance line in the sand, as this line has already been shown to have "no ontological basis" (Rosnow & Rosenthal, 1989, p. 1277).

Instead, we need to embrace the multitude of perspectives available to us and to begin to employ these diverse approaches. While NHST can still serve us well when properly utilized, it is important for researchers understand that different methods seek to answer different questions, and that we need to ensure that we are using the right method to answer a given question. When evaluating evidence in order to answer these questions we must be wary of looking for significant differences and focus instead to find meaningful differences. By combining these approaches we may be better able to qualify the strength of our evidence and discuss a more nuanced version of our data. Additionally, while all of these methods have drawbacks, when taken combined these methods can begin to overcome many of these limitations. For instance, given a large sample size, we would expect BF values to be very large and *p*-values to be very small, both indicating that the null model/hypothesis should not be supported. However, if we also have a PCC value of .30, we may decide that it is possible that this effect is very small and possibly negligible. This multifaceted approach may help to curb our enthusiasm over small or negligible "significant" effects that may not be practically meaningful and possibly may not replicate. Regardless if analyses agree or disagree on the presence of an effect, a researcher can investigate the size of the effect and discuss conclusions accordingly. Each methodology behaves slightly differently in given data environments, which might begin to highlight meaningful differences when discussed together.

Some may contest that all of these analyses are capable of being hacked, like *p*-values,

through researcher degrees of freedom, choice of priors, or pattern choice, among other actions (Simmons et al., 2011). Transparency throughout the research process is key to eliminating these issues, as α changes may only encourage bad research practices with the current incentive structure on publishing. Although we have the capability to share research across the world, research often still occurs behind closed doors. The Open Science Framework grants insight into research processes, allowing researchers to share their methodologies, code, design, and other important components of their projects. In addition to posting materials for projects, pre-registration of hypotheses and methodology will be an important facet in scientific accountability. Further, with increased transparency editors and other researchers can weigh the evidence presented according to their own beliefs.

Our key suggestion in this project is the redefinition of evidentiary value. The current focus on p-values has shown to be problematic, as many of the studies from the Open Science Collaboration (2015) do not replicate at $p < .05$ or $p < .005$ (Lakens et al., 2018). With the change in transparency mentioned above, publishing research with solid research designs and statistics, regardless of p-values, will allow for a broader range of evidence to become available. Publishing null findings is critical in replication and extension for discovering the limits and settings necessary for phenomena. Registered replications and reports will allow studies to be accepted prior to results being known, thus allowing researchers to focus on experimental design and hypotheses *apriori* instead of p-values *post hoc*. Reports should describe multiple indicators of evidence, such as effect sizes, confidence intervals, power analyses, Bayes Factors, and other descriptive statistics (Finkel, Eastwick, & Reis, 2015; Nosek & Lakens, 2014; van 't Veer, Giner-Sorolla, Van't Veer, & Giner-Sorolla, 2016).

A misunderstanding of statistical power still plagues psychological sciences (Bakker, Hartgerink, Wicherts, & Maas, 2016), and the effect of sample size, especially small ones, was shown here by comparing the criterion available in these analyses. Often, individual research labs may not have the means to adequately power a proposed study. Multilab studies and collaboration with other scientists is fundamental to alleviating these issues,

while encouraging interdisciplinary science. Collaboration increases our statistical abilities, as every researcher cannot be expected to be proficient in all methods and analyses, but teams of researchers can be assembled to cover a wider range of statistical skills to provide adequate estimates of evidence in their reports. We understand that there may be resistance to the implementation of multiple methodologies as these new methodologies take time and effort to learn. However, through the use of free programs (JASP, R, OOM, Shiny) and tutorials (YouTube, Coursera, http://www.statstools.com), we believe all researchers are capable of learning these analyses. We believe that through the expansion of our analytical knowledge and application of these new methodologies, we can begin to attenuate some of the strain currently placed on psychological science and to increase the strength of evidence in our discipline.

References

American Psychological Association. (2010). *Publication manual of the American Psychological Association* (6th ed., p. 272). American Psychological Association. Retrieved from http://www.apastyle.org/products/4200067.aspx

Bakker, M., Hartgerink, C. H. J., Wicherts, J. M., & Maas, H. L. J. van der. (2016). Researchers' intuitions about power in psychological research. *Psychological Science*, *27*(8), 1069–1077. doi:10.1177/0956797616647519

Bakker, M., Van Dijk, A., & Wicherts, J. M. (2012). The rules of the game called psychological science. *Perspectives on Psychological Science*, *7*(6), 543–554. doi:10.1177/1745691612459060

Bellhouse, D. R. (2004). The Reverend Thomas Bayes, FRS: A Biography to celebrate the tercentenary of his birth. *Statistical Science*, *19*(1), 3–43. doi:10.1214/088342304000000189

Benjamin, D. J., Berger, J. O., Johannesson, M., Nosek, B. A., Wagenmakers, E.-J., Berk, R., . . . Johnson, V. E. (2018). Redefine statistical significance. *Nature Human Behaviour*, *2*(1), 6–10. doi:10.1038/s41562-017-0189-z

Buchanan, E. M., Valentine, K. D., & Scofield, J. E. (2017). MOTE. Retrieved from https://github.com/doomlab/MOTE

Cohen, J. (1992). A power primer. *Psychological Bulletin*, *112*(1), 155–159. doi:10.1037/0033-2909.112.1.155

Cumming, G. (2008). Replication and p intervals. *Perspectives on Psychological Science*, *3*(4), 286–300. doi:10.1111/j.1745-6924.2008.00079.x

Cumming, G. (2014). The new statistics: Why and how. *Psychological Science*, *25*(1), 7–29. doi:10.1177/0956797613504966

Datta, G., & Ghosh, M. (1996). On the Invariance of Noninformative Priors. *The Annals of Statistics*, *24*(1), 141–159. doi:10.1214/aos/1033066203

De Laplace, P. S. (1774). Mémoire sur les suites récurro-récurrentes et sur leurs usages dans

la théorie des hasards. *Mém. Acad. R. Sci. Paris*, *6*(8), 353–371. Retrieved from

http://cerebro.cs.xu.edu/math/Sources/Laplace/recurro{_}recurrentes.pdf

Dienes, Z. (2008). *Understanding psychology as a science: an introduction to scientific and statistical inference.* Palgrave Macmillan.

Dienes, Z. (2014). Using Bayes to get the most out of non-significant results. *Frontiers in Psychology*, *5*(July), 1–17. doi:10.3389/fpsyg.2014.00781

Etz, A., & Vandekerckhove, J. (2016). A Bayesian perspective on the reproducibility project: Psychology. *PLoS ONE*, *11*(2), 1–12. doi:10.1371/journal.pone.0149794

Finkel, E. J., Eastwick, P. W., & Reis, H. T. (2015). Best research practices in psychology: Illustrating epistemological and pragmatic considerations with the case of relationship science. *Journal of Personality and Social Psychology*, *108*(2), 275–297. doi:10.1037/pspi0000007

Fisher, R. A. (1932). Inverse probability and the use of Likelihood. *Mathematical Proceedings of the Cambridge Philosophical Society*, *28*(3), 257–261. doi:10.1017/S0305004100010094

Gelman, A., Carlin, J. B., Stern, H. S., & Rubin, D. R. (2013). *Bayesian data analysis.* Chapman & Hall/CRC.

Genz, A., Bretz, F., Miwa, T., Mi, X., Leisch, F., Scheipl, F., & Hothorn, T. (2017). mvtnorm: Multivariate normal and t distributions. Retrieved from http://cran.r-project.org/package=mvtnorm

Gigerenzer, G. (2004). Mindless statistics. *The Journal of Socio-Economics*, *33*(5), 587–606. doi:10.1016/j.socec.2004.09.033

Gigerenzer, G., Krauss, S., & Vitouch, O. (2004). The null ritual: What you always wanted to know about significance testing but were afraid to ask. In *The sage handbook of quantitative methodology for the social sciences* (pp. 392–409). Thousand Oaks, CA: SAGE Publications, Inc. doi:10.4135/9781412986311.n21

Grice, J. W. (2011). *Observation Oriented Modeling : analysis of cause in the behavioral*

sciences (p. 242). Elsevier/Academic Press.

Grice, J. W. (2014). Observation Oriented Modeling: Preparing students for research in the 21st century. *Comprehensive Psychology*, *3*, 05.08.IT.3.3. doi:10.2466/05.08.IT.3.3

Grice, J. W., Barrett, P. T., Schlimgen, L. A., & Abramson, C. I. (2012). Toward a brighter future for psychology as an observation oriented science. *Behavioral Sciences*, *2*(4), 1–22. doi:10.3390/bs2010001

Grice, J. W., Craig, D. P. A., & Abramson, C. I. (2015). A simple and transparent alternative to repeated measures ANOVA. *SAGE Open*, *5*(3), 2158244015604192. doi:10.1177/2158244015604192

Haaf, J., & Rouder, J. N. (2017). *Developing constraint in bayesian mixed models.* doi:10.17605/OSF.IO/KTJNQ

Ioannidis, J. P. A. (2005). Why most published research findings are false. *PLoS Medicine*, *2*(8), e124. doi:10.1371/journal.pmed.0020124

JASP Team. (2017). JASP. Retrieved from https://jasp-stats.org/

Klugkist, I., & Hoijtink, H. (2007). The Bayes factor for inequality and about equality constrained models. *Computational Statistics & Data Analysis*, *51*(12), 6367–6379. doi:10.1016/j.csda.2007.01.024

Kruschke, J. K. (2014). *Doing Bayesian data analysis: A tutorial with R, JAGS, and Stan* (2nd ed.). Academic Press.

Lakens, D. (2013). Calculating and reporting effect sizes to facilitate cumulative science: A practical primer for t-tests and ANOVAs. *Frontiers in Psychology*, *4*. doi:10.3389/fpsyg.2013.00863

Lakens, D. (2017). Equivalence tests. *Social Psychological and Personality Science*, *8*(4), 355–362. doi:10.1177/1948550617697177

Lakens, D., Adolfi, F. G., Albers, C. J., Anvari, F., Apps, M. A. J., Argamon, S. E., . . . Zwaan, R. A. (2018). Justify your alpha. *Nature Human Behaviour*, *2*(3), 168–171.

doi:10.1038/s41562-018-0311-x

Lawrence, M. A. (2017). ez: Easy analysis and visualization of factorial experiments.

Retrieved from http://cran.r-project.org/package=ez

Lehmann, E. L. (1993). The Fisher, Neyman-Pearson theories of testing hypotheses: One

theory or two? *Journal of the American Statistical Association*, *88*(424), 1242–1249.

doi:10.1080/01621459.1993.10476404

Lehmann, E. L. (2011). *Fisher, Neyman, and the creation of classical statistics*. New York,

NY: Springer.

Lindsay, D. S. (2015). Replication in Psychological Science. *Psychological Science*, *26*(12),

1827–1832. doi:10.1177/0956797615616374

Maxwell, S. E., & Delaney, H. D. (2004). *Designing experiments and analyzing data: A*

model comparison perspective (2nd ed.). Mahwah, NJ: Lawrence Erlbaum Associates.

Maxwell, S. E., Lau, M. Y., & Howard, G. S. (2015). Is psychology suffering from a

replication crisis? What does "failure to replicate" really mean? *American*

Psychologist, *70*(6), 487–498. doi:10.1037/a0039400

Morey, R. D. (2015). On verbal categories for the interpretation of Bayes factors. Retrieved

from http:

//bayesfactor.blogspot.com/2015/01/on-verbal-categories-for-interpretation.html

Morey, R. D., & Rouder, J. N. (2015). BayesFactor: Computation of Bayes Factors for

common designs. Retrieved from https://cran.r-project.org/package=BayesFactor

Nosek, B. A., & Lakens, D. (2014). Registered reports. *Social Psychology*, *45*(3), 137–141.

doi:10.1027/1864-9335/a000192

Nosek, B. A., Spies, J. R., & Motyl, M. (2012). Scientific utopia. *Perspectives on*

Psychological Science, *7*(6), 615–631. doi:10.1177/1745691612459058

Open Science Collaboration. (2015). Estimating the reproducibility of psychological science.

Science, *349*(6251), aac4716–aac4716. doi:10.1126/science.aac4716

Pericchi, L., & Pereira, C. (2016). Adaptive significance levels using optimal decision rules:

Balancing by weighting the error probabilities. *Brazilian Journal of Probability and Statistics*, *30*(1), 70–90. doi:10.1214/14-BJPS257

Press, S. J. (2002). *Subjective and objective Bayesian statistics* (2nd ed.). Hoboken, NJ, USA: John Wiley & Sons, Inc. doi:10.1002/9780470317105

Rosnow, R. L., & Rosenthal, R. (1989). Statistical procedures and the justification of knowledge in psychological science. *American Psychologist*, *44*(10), 1276–1284. doi:10.1037/0003-066X.44.10.1276

Rouder, J. N., Morey, R. D., Speckman, P. L., & Province, J. M. (2012). Default Bayes factors for ANOVA designs. *Journal of Mathematical Psychology*, *56*(5), 356–374. doi:10.1016/j.jmp.2012.08.001

Rouder, J. N., Speckman, P. L., Sun, D., Morey, R. D., & Iverson, G. (2009). Bayesian t tests for accepting and rejecting the null hypothesis. *Psychonomic Bulletin & Review*, *16*(2), 225–237. doi:10.3758/PBR.16.2.225

Sauer, S., & Luebke, K. (2016, January). Observation Oriented Modeling revised from a statistical point of view. doi:10.17605/OSF.IO/3J4XR

Sellke, T., Bayarri, M. J., & Berger, J. O. (2001). Calibration of p values for testing precise null hypotheses. *American Statistician*, *55*(1), 62–71. doi:10.1198/000313001300339950

Simmons, J. P., Nelson, L. D., & Simonsohn, U. (2011). False-positive psychology: Undisclosed flexibility in data collection and analysis allows presenting anything as significant. *Psychological Science*, *22*(11), 1359–1366. doi:10.1177/0956797611417632

Tabachnick, B. G., & Fidell, L. S. (2012). *Using multivariate statistics* (Sixth.). Boston, MA: Pearson.

Valentine, K. D., & Buchanan, E. M. (2013). JAM-boree: An application of observation oriented modelling to judgements of associative memory. *Journal of Cognitive Psychology*, *25*(4), 400–422. doi:10.1080/20445911.2013.775120

Van Elk, M., Matzke, D., Gronau, Q. F., Guan, M., Vandekerckhove, J., & Wagenmakers,

E.-J. (2015). Meta-analyses are no substitute for registered replications: A skeptical perspective on religious priming. *Frontiers in Psychology*, *6*, 1365. doi:10.3389/fpsyg.2015.01365

van 't Veer, A. E., Giner-Sorolla, R., Van't Veer, A. E., & Giner-Sorolla, R. (2016). Pre-registration in social psychology—A discussion and suggested template. *Journal of Experimental Social Psychology*, *67*, 2–12. doi:10.1016/j.jesp.2016.03.004

Wagenmakers, E.-J. (2007). A practical solution to the pervasive problems of p values. *Psychonomic Bulletin & Review*, *14*(5), 779–804. doi:10.3758/BF03194105

Wasserstein, R. L., & Lazar, N. A. (2016). The ASA's statement on p -values: Context, process, and purpose. *The American Statistician*, *70*(2), 129–133. doi:10.1080/00031305.2016.1154108

Wetzels, R., Matzke, D., Lee, M. D., Rouder, J. N., Iverson, G. J., & Wagenmakers, E.-J. (2011). Statistical evidence in experimental psychology. *Perspectives on Psychological Science*, *6*(3), 291–298. doi:10.1177/1745691611406923
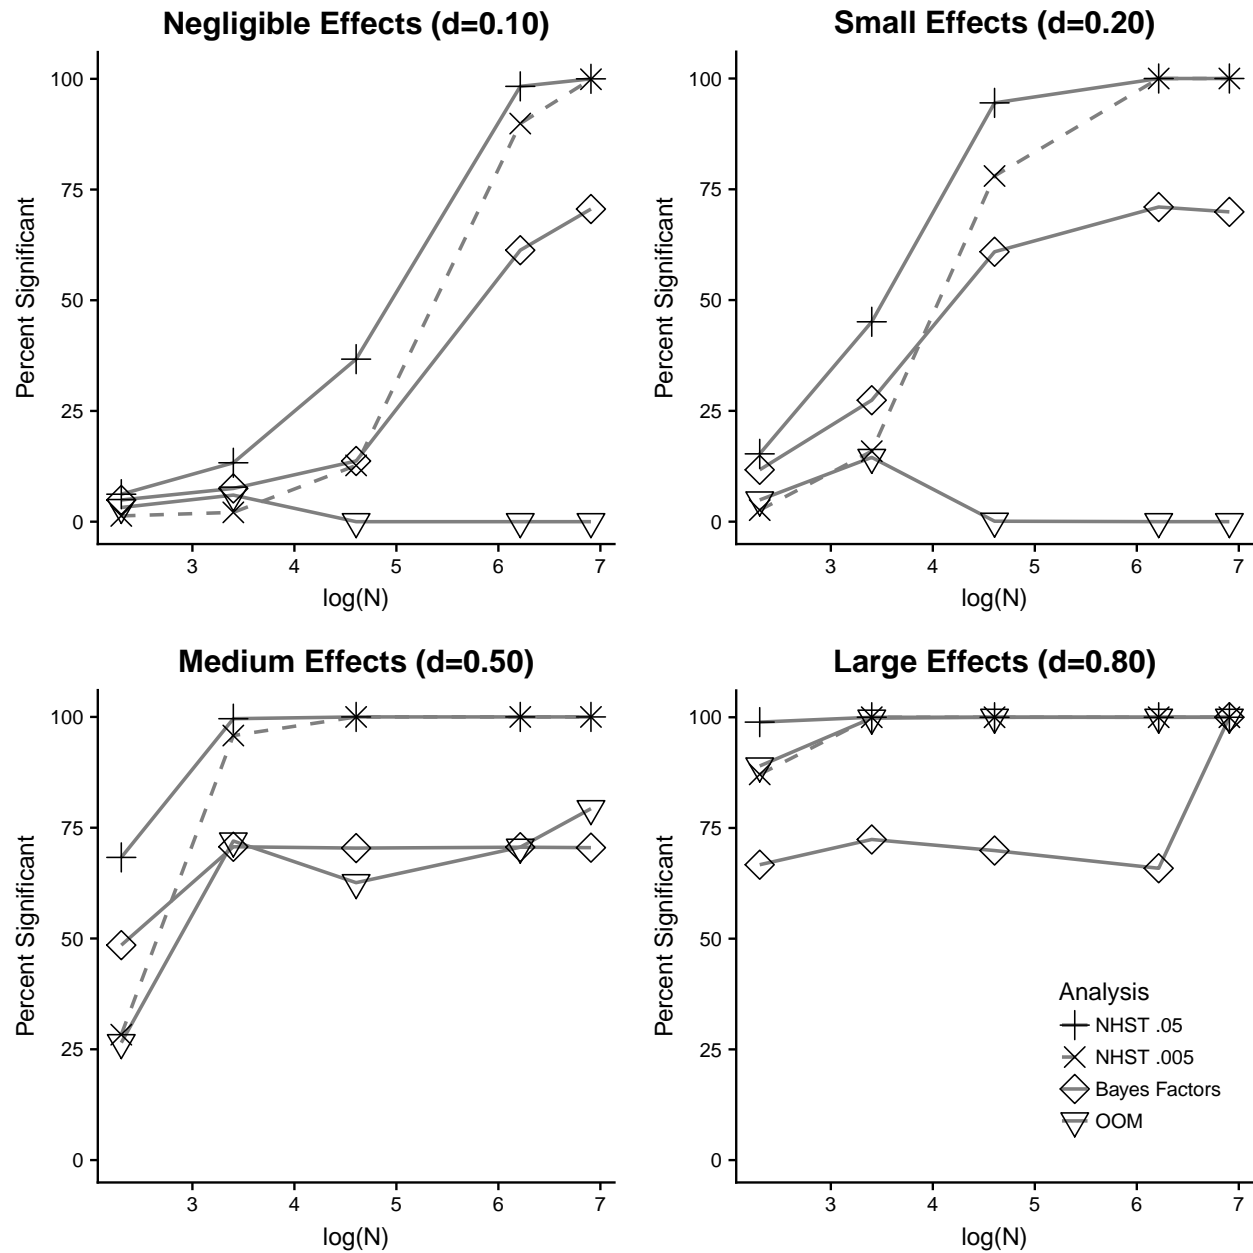
Figure 1. For NHST analyses only, percent of significant estimates at $p < .05$ (solid) and $p <$ .005 (dashed) for each analysis given effect size and sample size.

Figure 2. Percent of non-significant effects for each analysis given effect size and sample size.
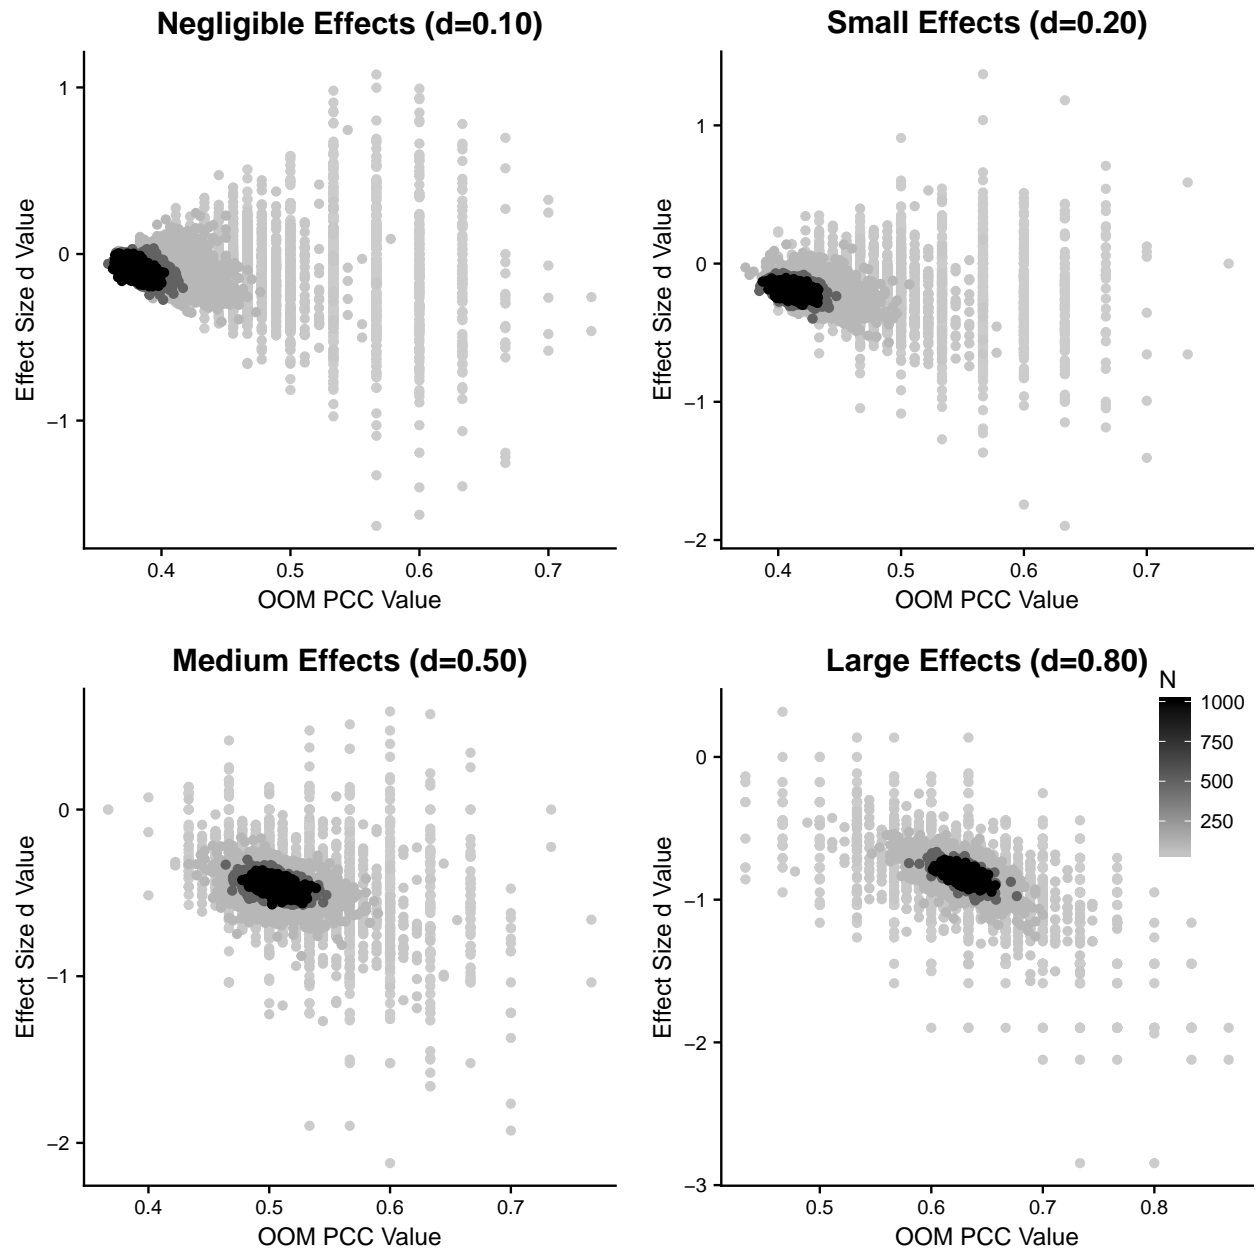
Figure 3. WHAT DO YOU WANT THIS TO SAY

Figure 4. Percent of agreement across all analyses given effect size and sample size for omnibus tests.

Figure 5. Percent of agreement across each analysis given effect size and sample size *posthoc* tests with $p < .05$ (solid) and $p < .005$ (dashed). Note that this graph only compares the NHST and BF conclusions.
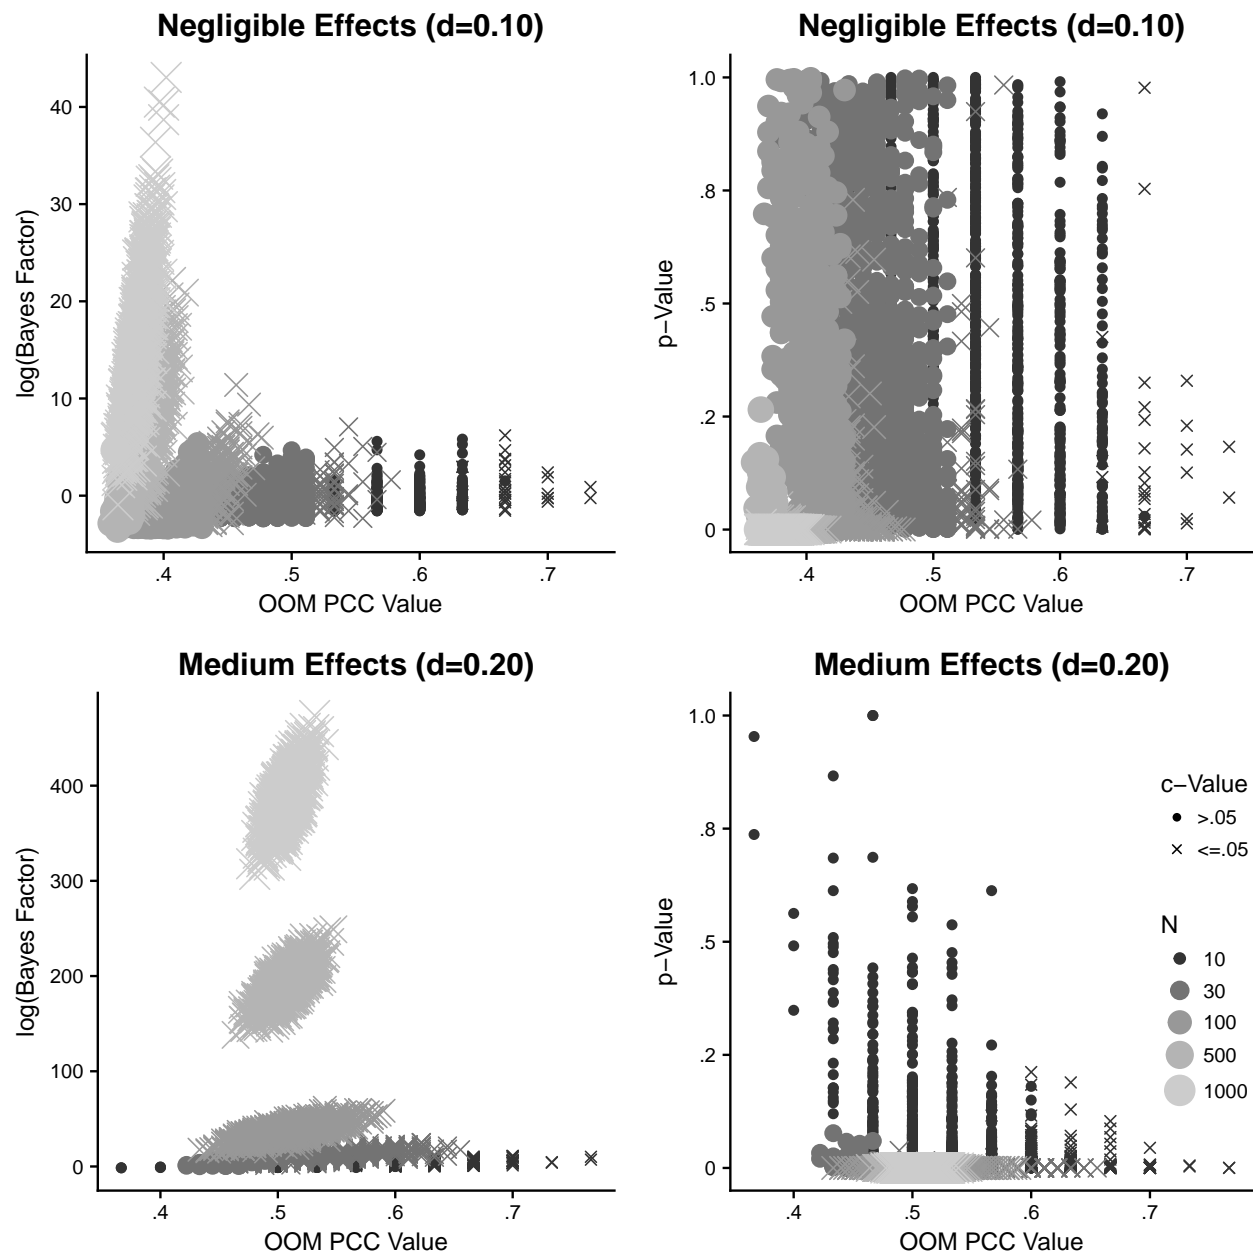
Figure 6. PCC and c-values plotted against *p* and BF values for negligible and medium effect size conditions. Xs indicate simulations with *c*-values < .05,which were binned as significant if they were found in conjunction with PCC values over .50. Point size and color indicates sample size wherein larger samples are lighter colors. Please note that the key for the entire set of graphics is in the bottom right hand corner.